

```

In [ ]: from __future__ import absolute_import, division, print_function

from options import LiteMonoOptions
from trainer import Trainer

from __future__ import absolute_import, division, print_function

import time
import torch.optim as optim
from torch.utils.data import DataLoader
from tensorboardX import SummaryWriter

import json

from utils import *
from kitti_utils import *
from layers import *

import datasets
import networks
from linear_warmup_cosine_annealing_warm_restarts_weight_decay import ChainedScheduler
import albumentations as A
from albumentations.pytorch import ToTensorV2

# torch.backends.cudnn.benchmark = True

import numpy as np
import pandas as pd
import torch
import torch.nn as nn
from torch.cuda.amp import GradScaler, autocast
import matplotlib.pyplot as plt
from tqdm.auto import tqdm
from pathlib import Path
import cv2 as cv
from PIL import Image
import segmentation_models_pytorch as smp
from sklearn.model_selection import train_test_split
import albumentations as A
from albumentations.pytorch import ToTensorV2
from torchmetrics.image import StructuralSimilarityIndexMeasure as SSIM
from torchmetrics.regression import MeanSquaredError as MSE
from torchmetrics.collections import MetricCollection
import gc
from torchvision.transforms import Normalize

def time_sync():
    # PyTorch-accurate time
    if torch.cuda.is_available():
        torch.cuda.synchronize()
    return time.time()

sample_tfms = [
    A.HorizontalFlip(),
    A.GaussNoise(p=0.2),
    A.OneOf([
        A.MotionBlur(p=.3),
        A.MedianBlur(blur_limit=3, p=0.3),
        A.Blur(blur_limit=3, p=0.5),
    ], p=0.3),
    A.RGBShift(),
    A.RandomBrightnessContrast(),
    A.RandomResizedCrop(384,384),
    A.ColorJitter(),
    A.ShiftScaleRotate(shift_limit=0.1, scale_limit=0.3, rotate_limit=45, p=0.5),
    A.HueSaturationValue(p=0.3),
]
train_tfms = A.Compose([
    *sample_tfms,
    A.Resize(224,224),
    A.Normalize(always_apply=True),
    ToTensorV2()
])
valid_tfms = A.Compose([

```

```

A.Resize(224,224),
A.Normalize(always_apply=True),
ToTensorV2()
])

class Dataset:
    def __init__(self,df,tfms):
        self.df = df
        self.tfms=tfms
    def open_im(self,p,gray=False):
        im = cv.imread(str(p))
        im = cv.cvtColor(im,cv.COLOR_BGR2GRAY if gray else cv.COLOR_BGR2RGB)
        return im

    def __len__(self,):
        return len(self.df)

    def __getitem__(self,idx):
        s = self.df.iloc[idx,:]
        im, dp = s[0],s[1]
        im, dp = self.open_im(im), self.open_im(dp,True)
        augs = self.tfms(image=im,mask=dp)
        im, dp = augs['image'], augs['mask'] / 255.
        return im, dp.unsqueeze(0)

train_csv = Path('./nyu-depth-v2/nyu_data/data/nyu2_train.csv')
train_ims_path = Path('./nyu-depth-v2/nyu_data/data/nyu2_train')
base_path = Path('./nyu-depth-v2/nyu_data')

df = pd.read_csv(train_csv,header=None)
df[0] = df[0].map(lambda x:base_path/x)
df[1] = df[1].map(lambda x:base_path/x)
df.head()

train_df, val_df = train_test_split(df,test_size=0.1,shuffle=True)
val_df, test_df = train_test_split(val_df, test_size=0.1,shuffle=True)
train_df.reset_index(drop=True,inplace=True)
val_df.reset_index(drop=True,inplace=True)
test_df.reset_index(drop=True,inplace=True)
len(train_df),len(val_df), len(test_df)

train_ds = Dataset(train_df,train_tfms)
val_ds = Dataset(val_df,valid_tfms)
test_ds = Dataset(test_df, valid_tfms)
len(train_ds), len(val_ds), len(test_ds)

fn_loss = nn.MSELoss()

class Trainer:
    def __init__(self, options):
        self.opt = options
        self.log_path = os.path.join(self.opt.log_dir, self.opt.model_name)

        # checking height and width are multiples of 32
        assert self.opt.height % 32 == 0, "'height' must be a multiple of 32"
        assert self.opt.width % 32 == 0, "'width' must be a multiple of 32"

        self.models = {}
        self.models_pose = {}
        self.parameters_to_train = []
        self.parameters_to_train_pose = []

        self.device = torch.device("cpu" if self.opt.no_cuda else "cuda")
        self.profile = self.opt.profile

        self.num_scales = len(self.opt.scales)
        self.frame_ids = len(self.opt.frame_ids)
        self.num_pose_frames = 2 if self.opt.pose_model_input == "pairs" else self.num_input_frames

        assert self.opt.frame_ids[0] == 0, "frame_ids must start with 0"

        self.use_pose_net = not (self.opt.use_stereo and self.opt.frame_ids == [0])

        if self.opt.use_stereo:
            self.opt.frame_ids.append("s")

        self.models["encoder"] = networks.LiteMono(model=self.opt.model,
                                                    drop_path_rate=self.opt.drop_path,
                                                    width=self.opt.width, height=self.opt.height)

```

```

self.models["encoder"].to(self.device)
self.parameters_to_train += list(self.models["encoder"].parameters())

self.models["depth"] = networks.DepthDecoder(self.models["encoder"].num_ch_enc,
                                              self.opt.scales)
self.models["depth"].to(self.device)
self.parameters_to_train += list(self.models["depth"].parameters())

if self.use_pose_net:
    if self.opt.pose_model_type == "separate_resnet":
        self.models_pose["pose_encoder"] = networks.ResnetEncoder(
            self.opt.num_layers,
            self.opt.weights_init == "pretrained",
            num_input_images=self.num_pose_frames)

        self.models_pose["pose_encoder"].to(self.device)
        self.parameters_to_train_pose += list(self.models_pose["pose_encoder"].parameters())

        self.models_pose["pose"] = networks.PoseDecoder(
            self.models_pose["pose_encoder"].num_ch_enc,
            num_input_features=1,
            num_frames_to_predict_for=2)

    elif self.opt.pose_model_type == "shared":
        self.models_pose["pose"] = networks.PoseDecoder(
            self.models["encoder"].num_ch_enc, self.num_pose_frames)

    elif self.opt.pose_model_type == "posecnn":
        self.models_pose["pose"] = networks.PoseCNN(
            self.num_input_frames if self.opt.pose_model_input == "all" else 2)

        self.models_pose["pose"].to(self.device)
        self.parameters_to_train_pose += list(self.models_pose["pose"].parameters())

if self.opt.predictive_mask:
    assert self.opt.disable_automasking, \
        "When using predictive_mask, please disable automasking with --disable_automasking"

    # Our implementation of the predictive masking baseline has the the same architecture
    # as our depth decoder. We predict a separate mask for each source frame.
    self.models["predictive_mask"] = networks.DepthDecoder(
        self.models["encoder"].num_ch_enc, self.opt.scales,
        num_output_channels=(len(self.opt.frame_ids) - 1))
    self.models["predictive_mask"].to(self.device)
    self.parameters_to_train += list(self.models["predictive_mask"].parameters())

self.model_optimizer = optim.AdamW(self.parameters_to_train, self.opt.lr[0], weight_decay=)
if self.use_pose_net:
    self.model_pose_optimizer = optim.AdamW(self.parameters_to_train_pose, self.opt.lr[3],

self.model_lr_scheduler = ChainedScheduler(
    self.model_optimizer,
    T_0=int(self.opt.lr[2]),
    T_mul=1,
    eta_min=self.opt.lr[1],
    last_epoch=-1,
    max_lr=self.opt.lr[0],
    warmup_steps=0,
    gamma=0.9
)
self.model_pose_lr_scheduler = ChainedScheduler(
    self.model_pose_optimizer,
    T_0=int(self.opt.lr[5]),
    T_mul=1,
    eta_min=self.opt.lr[4],
    last_epoch=-1,
    max_lr=self.opt.lr[3],
    warmup_steps=0,
    gamma=0.9
)

if self.opt.load_weights_folder is not None:
    self.load_model()

if self.opt.mypretrain is not None:
    self.load_pretrain()

# data

```

```

datasets_dict = {"kitti": datasets.KITTIRAWDataset,
                  "kitti_odom": datasets.KITTIOdomDataset}
self.dataset = datasets_dict[self.opt.dataset]

fpath = os.path.join(par_dir, "splits", self.opt.split, "{}_files.txt")

train_filenames = readlines(fpath.format("train"))
val_filenames = readlines(fpath.format("val"))
img_ext = '.png' if self.opt.png else '.jpg'

num_train_samples = len(train_filenames)
self.num_total_steps = num_train_samples // self.opt.batch_size * self.opt.num_epochs
#####
train_dataset = train_ds
#self.dataset(
#    # self.opt.data_path, train_filenames, self.opt.height, self.opt.width,
#    # self.opt.frame_ids, 4, is_train=True, img_ext=img_ext)
self.train_loader = DataLoader(
    train_dataset, self.opt.batch_size, True,
    num_workers=self.opt.num_workers, pin_memory=True, drop_last=True)
val_dataset = val_ds
# self.dataset(
#     self.opt.data_path, val_filenames, self.opt.height, self.opt.width,
#     self.opt.frame_ids, 4, is_train=False, img_ext=img_ext)
self.val_loader = DataLoader(
    val_dataset, self.opt.batch_size, True,
    num_workers=self.opt.num_workers, pin_memory=True, drop_last=True)
self.val_iter = iter(self.val_loader)

self.writers = {}
for mode in ["train", "val"]:
    self.writers[mode] = SummaryWriter(os.path.join(self.log_path, mode))

if not self.opt.no_ssim:
    self.ssim = SSIM()
    self.ssim.to(self.device)

self.backproject_depth = {}
self.project_3d = {}
for scale in self.opt.scales:
    h = self.opt.height // (2 ** scale)
    w = self.opt.width // (2 ** scale)

    self.backproject_depth[scale] = BackprojectDepth(self.opt.batch_size, h, w)
    self.backproject_depth[scale].to(self.device)

    self.project_3d[scale] = Project3D(self.opt.batch_size, h, w)
    self.project_3d[scale].to(self.device)

self.depth_metric_names = [
    "de/abs_rel", "de/sq_rel", "de/rms", "de/log_rms", "da/a1", "da/a2", "da/a3"]

print("There are {:d} training items and {:d} validation items\n".format(
    len(train_dataset), len(val_dataset)))

self.save_opts()

def set_train(self):
    """Convert all models to training mode
    """
    for m in self.models.values():
        m.train()

def set_eval(self):
    """Convert all models to testing/evaluation mode
    """
    for m in self.models.values():
        m.eval()

def train(self):
    """Run the entire training pipeline
    """
    self.epoch = 0
    self.step = 0
    self.start_time = time.time()
    for self.epoch in range(self.opt.num_epochs):
        self.run_epoch()
        if (self.epoch + 1) % self.opt.save_frequency == 0:
            self.save_model()

```

```

def run_epoch(self):
    """Run a single epoch of training and validation
    """

    print("Training")
    self.set_train()

    self.model_lr_scheduler.step()
    if self.use_pose_net:
        self.model_pose_lr_scheduler.step()

    for batch_idx, inputs in enumerate(self.train_loader):

        before_op_time = time.time()

        outputs, losses = self.process_batch(inputs)

        # self.model_optimizer.zero_grad()
        # if self.use_pose_net:
        #     self.model_pose_optimizer.zero_grad()
        losses.backward()
        self.model_optimizer.step()
        # if self.use_pose_net:
        #     self.model_pose_optimizer.step()

        duration = time.time() - before_op_time

        # log less frequently after the first 2000 steps to save time & disk space
        early_phase = batch_idx % self.opt.log_frequency == 0 and self.step < 20000
        late_phase = self.step % 2000 == 0

        if early_phase or late_phase:
            self.log_time(batch_idx, duration, losses.cpu().data)

            # if "depth_gt" in inputs:
            #     self.compute_depth_losses(inputs, outputs, losses)

            # self.log("train", inputs, outputs, losses)
            self.val()

        self.step += 1

def process_batch(self, inputs):
    """Pass a minibatch through the network and generate images and losses
    """
    losses = 0.
    # print(inputs[0].shape, inputs[1].shape, len(inputs))
    for key, ipt in enumerate(inputs):
        inputs[key] = ipt.to(self.device)

    if self.opt.pose_model_type == "shared":
        # If we are using a shared encoder for both depth and pose (as advocated
        # in monodepthv1), then all images are fed separately through the depth encoder.
        all_color_aug = torch.cat([inputs[("color_aug", i, 0)] for i in self.opt.frame_ids])
        all_features = self.models["encoder"](all_color_aug)
        all_features = [torch.split(f, self.opt.batch_size) for f in all_features]

        features = {}
        for i, k in enumerate(self.opt.frame_ids):
            features[k] = [f[i] for f in all_features]

        outputs = self.models["depth"](features[0])
    else:
        # Otherwise, we only feed the image with frame_id 0 through the depth encoder

    features = self.models["encoder"](inputs[0])[("color_aug", 0, 0)]

    outputs = self.models["depth"](features)

    # if self.opt.predictive_mask:
    #     outputs["predictive_mask"] = self.models["predictive_mask"](features)

    # if self.use_pose_net:
    #     outputs.update(self.predict_poses(inputs, features))

    # self.generate_images_pred(inputs[0], outputs)
    # print(outputs.keys())
    # print(list(outputs.values())[-1].shape)
    # print(list(outputs.values())[-1])
    # losses = self.compute_losses(inputs[1], list(outputs.values())[-1])

```

```

gts = inputs[1]
prds = list(outputs.values())[-1]
for prd, gt in zip(gts, prds):
    loss = fn_loss(gts, prds)
    losses += loss
# losses = nn.MSELoss(inputs[1], list(outputs.values())[-1])

return outputs, losses

def predict_poses(self, inputs, features):
    """Predict poses between input frames for monocular sequences.
    """
    outputs = {}
    if self.num_pose_frames == 2:
        # In this setting, we compute the pose to each source frame via a
        # separate forward pass through the pose network.

        # select what features the pose network takes as input
        if self.opt.pose_model_type == "shared":
            pose_feats = {f_i: features[f_i] for f_i in self.opt.frame_ids}
        else:
            pose_feats = {f_i: inputs["color_aug", f_i, 0] for f_i in self.opt.frame_ids}

        for f_i in self.opt.frame_ids[1:]:
            if f_i != "s":
                # To maintain ordering we always pass frames in temporal order
                if f_i < 0:
                    pose_inputs = [pose_feats[f_i], pose_feats[0]]
                else:
                    pose_inputs = [pose_feats[0], pose_feats[f_i]]

                if self.opt.pose_model_type == "separate_resnet":
                    pose_inputs = [self.models_pose["pose_encoder"](torch.cat(pose_inputs, 1))]
                elif self.opt.pose_model_type == "posecnn":
                    pose_inputs = torch.cat(pose_inputs, 1)

                axisangle, translation = self.models_pose["pose"](pose_inputs)
                outputs[("axisangle", 0, f_i)] = axisangle
                outputs[("translation", 0, f_i)] = translation

                # Invert the matrix if the frame id is negative
                outputs[("cam_T_cam", 0, f_i)] = transformation_from_parameters(
                    axisangle[:, 0], translation[:, 0], invert=(f_i < 0))

            else:
                # Here we input all frames to the pose net (and predict all poses) together
                if self.opt.pose_model_type in ["separate_resnet", "posecnn"]:
                    pose_inputs = torch.cat(
                        [inputs[("color_aug", i, 0)] for i in self.opt.frame_ids if i != "s"], 1)

                    if self.opt.pose_model_type == "separate_resnet":
                        pose_inputs = [self.models["pose_encoder"](pose_inputs)]

                elif self.opt.pose_model_type == "shared":
                    pose_inputs = [features[i] for i in self.opt.frame_ids if i != "s"]

                axisangle, translation = self.models_pose["pose"](pose_inputs)

                for i, f_i in enumerate(self.opt.frame_ids[1:]):
                    if f_i != "s":
                        outputs[("axisangle", 0, f_i)] = axisangle
                        outputs[("translation", 0, f_i)] = translation
                        outputs[("cam_T_cam", 0, f_i)] = transformation_from_parameters(
                            axisangle[:, i], translation[:, i])

            return outputs

def val(self):
    """Validate the model on a single minibatch
    """
    self.set_eval()
    try:
        inputs = self.val_iter.next()
    except StopIteration:
        self.val_iter = iter(self.val_loader)
        inputs = self.val_iter.next()

    with torch.no_grad():
        outputs, losses = self.process_batch(inputs)

```

```

        if "depth_gt" in inputs:
            self.compute_depth_losses(inputs, outputs, losses)

        # self.log("val", inputs, outputs, losses)
        del inputs, outputs, losses

    self.set_train()

def generate_images_pred(self, inputs, outputs):
    """Generate the warped (reprojected) color images for a minibatch.
    Generated images are saved into the `outputs` dictionary.
    """
    for scale in self.opt.scales:
        disp = outputs[("disp", scale)]
        if self.opt.v1_multiscale:
            source_scale = scale
        else:
            disp = F.interpolate(
                disp, [self.opt.height, self.opt.width], mode="bilinear", align_corners=False)
            source_scale = 0

        _, depth = disp_to_depth(disp, self.opt.min_depth, self.opt.max_depth)

        outputs[("depth", 0, scale)] = depth

        for i, frame_id in enumerate(self.opt.frame_ids[1:]):

            if frame_id == "s":
                T = inputs["stereo_T"]
            else:
                T = outputs[("cam_T_cam", 0, frame_id)]

            # from the authors of https://arxiv.org/abs/1712.00175
            if self.opt.pose_model_type == "posecnn":

                axisangle = outputs[("axisangle", 0, frame_id)]
                translation = outputs[("translation", 0, frame_id)]

                inv_depth = 1 / depth
                mean_inv_depth = inv_depth.mean(3, True).mean(2, True)

                T = transformation_from_parameters(
                    axisangle[:, 0], translation[:, 0] * mean_inv_depth[:, 0], frame_id < 0)

            cam_points = self.backproject_depth[source_scale](
                depth, inputs[("inv_K", source_scale)])
            pix_coords = self.project_3d[source_scale](
                cam_points, inputs[("K", source_scale)], T)

            outputs[("sample", frame_id, scale)] = pix_coords

            outputs[("color", frame_id, scale)] = F.grid_sample(
                inputs[("color", frame_id, source_scale)],
                outputs[("sample", frame_id, scale)],
                padding_mode="border", align_corners=True)

            if not self.opt.disable_automasking:
                outputs[("color_identity", frame_id, scale)] = \
                    inputs[("color", frame_id, source_scale)]

def compute_reprojection_loss(self, pred, target):
    """Computes reprojection loss between a batch of predicted and target images
    """
    abs_diff = torch.abs(target - pred)
    l1_loss = abs_diff.mean(1, True)

    if self.opt.no_ssim:
        reprojection_loss = l1_loss
    else:
        ssim_loss = self.ssim(pred, target).mean(1, True)
        reprojection_loss = 0.85 * ssim_loss + 0.15 * l1_loss

    return reprojection_loss

def compute_losses(self, inputs, outputs):
    """Compute the reprojection and smoothness losses for a minibatch
    """

    losses = {}
    total_loss = 0

```

```

for scale in self.opt.scales:
    loss = 0
    reprojection_losses = []

    if self.opt.v1_multiscale:
        source_scale = scale
    else:
        source_scale = 0

    disp = outputs[("disp", scale)]
    color = inputs[("color", 0, scale)]
    target = inputs[("color", 0, source_scale)]

    for frame_id in self.opt.frame_ids[1:]:
        pred = outputs[("color", frame_id, scale)]
        reprojection_losses.append(self.compute_reprojection_loss(pred, target))

    reprojection_losses = torch.cat(reprojection_losses, 1)

    if not self.opt.disable_automasking:
        identity_reprojection_losses = []
        for frame_id in self.opt.frame_ids[1:]:
            pred = inputs[("color", frame_id, source_scale)]
            identity_reprojection_losses.append(
                self.compute_reprojection_loss(pred, target))

        identity_reprojection_losses = torch.cat(identity_reprojection_losses, 1)

        if self.opt.avg_reprojection:
            identity_reprojection_loss = identity_reprojection_losses.mean(1, keepdim=True)
        else:
            # save both images, and do min all at once below
            identity_reprojection_loss = identity_reprojection_losses

    elif self.opt.predictive_mask:
        # use the predicted mask
        mask = outputs["predictive_mask"][("disp", scale)]
        if not self.opt.v1_multiscale:
            mask = F.interpolate(
                mask, [self.opt.height, self.opt.width],
                mode="bilinear", align_corners=False)

        reprojection_losses *= mask

        # add a loss pushing mask to 1 (using nn.BCELoss for stability)
        weighting_loss = 0.2 * nn.BCELoss()(mask, torch.ones(mask.shape).cuda())
        loss += weighting_loss.mean()

    if self.opt.avg_reprojection:
        reprojection_loss = reprojection_losses.mean(1, keepdim=True)
    else:
        reprojection_loss = reprojection_losses

    if not self.opt.disable_automasking:
        # add random numbers to break ties
        identity_reprojection_loss += torch.randn(
            identity_reprojection_loss.shape, device=self.device) * 0.00001

    combined = torch.cat((identity_reprojection_loss, reprojection_loss), dim=1)
    else:
        combined = reprojection_loss

    if combined.shape[1] == 1:
        to_optimise = combined
    else:
        to_optimise, idxs = torch.min(combined, dim=1)

    if not self.opt.disable_automasking:
        outputs["identity_selection/{}".format(scale)] = (
            idxs > identity_reprojection_loss.shape[1] - 1).float()

    loss += to_optimise.mean()

    mean_disp = disp.mean(2, True).mean(3, True)
    norm_disp = disp / (mean_disp + 1e-7)
    smooth_loss = get_smooth_loss(norm_disp, color)

    loss += self.opt.disparity_smoothness * smooth_loss / (2 ** scale)
    total_loss += loss

```



```

        losses["loss/{}".format(scale)] = loss

    total_loss /= self.num_scales
    losses["loss"] = total_loss
    return losses

def compute_depth_losses(self, inputs, outputs, losses):
    """Compute depth metrics, to allow monitoring during training

    This isn't particularly accurate as it averages over the entire batch,
    so is only used to give an indication of validation performance
    """
    depth_pred = outputs[("depth", 0, 0)]
    depth_pred = torch.clamp(F.interpolate(
        depth_pred, [375, 1242], mode="bilinear", align_corners=False), 1e-3, 80)
    depth_pred = depth_pred.detach()

    depth_gt = inputs["depth_gt"]
    mask = depth_gt > 0

    # garg/eigen crop
    crop_mask = torch.zeros_like(mask)
    crop_mask[:, :, 153:371, 44:1197] = 1
    mask = mask * crop_mask

    depth_gt = depth_gt[mask]
    depth_pred = depth_pred[mask]
    depth_pred *= torch.median(depth_gt) / torch.median(depth_pred)

    depth_pred = torch.clamp(depth_pred, min=1e-3, max=80)

    depth_errors = compute_depth_errors(depth_gt, depth_pred)

    for i, metric in enumerate(self.depth_metric_names):
        losses[metric] = np.array(depth_errors[i].cpu())

def log_time(self, batch_idx, duration, loss):
    """Print a logging statement to the terminal
    """
    samples_per_sec = self.opt.batch_size / duration
    time_sofar = time.time() - self.start_time
    training_time_left = (
        self.num_total_steps / self.step - 1.0) * time_sofar if self.step > 0 else 0
    print_string = "epoch {:>3} | lr {:.6f} | lr_p {:.6f} | batch {:>6} | examples/s: {:.1f}" -
        " | loss: {:.5f} | time elapsed: {} | time left: {}"
    print(print_string.format(self.epoch, self.model_optimizer.state_dict()['param_groups'][0]
        self.model_pose_optimizer.state_dict()['param_groups'][0]['lr'],
        batch_idx, samples_per_sec, loss,
        sec_to_hm_str(time_sofar), sec_to_hm_str(training_time_left)))

def log(self, mode, inputs, outputs, losses):
    """Write an event to the tensorboard events file
    """
    writer = self.writers[mode]
    for l, v in losses.items():
        writer.add_scalar("{} {}".format(l, v), self.step)

    for j in range(min(4, self.opt.batch_size)): # write a maximum of four images
        for s in self.opt.scales:
            for frame_id in self.opt.frame_ids:
                writer.add_image(
                    "color_{}_{}_{}".format(frame_id, s, j),
                    inputs[("color", frame_id, s)][j].data, self.step)
                if s == 0 and frame_id != 0:
                    writer.add_image(
                        "color_pred_{}_{}_{}".format(frame_id, s, j),
                        outputs[("color", frame_id, s)][j].data, self.step)

            writer.add_image(
                "disp_{}_{}".format(s, j),
                normalize_image(outputs[("disp", s)][j]), self.step)

            if self.opt.predictive_mask:
                for f_idx, frame_id in enumerate(self.opt.frame_ids[1:]):
                    writer.add_image(
                        "predictive_mask_{}_{}_{}".format(frame_id, s, j),
                        outputs["predictive_mask"][("disp", s)][j, f_idx][None, ...],
                        self.step)

            elif not self.opt.disable_automasking:

```

```

        writer.add_image(
            "automask_{}/{}".format(s, j),
            outputs["identity_selection/{}".format(s)][j][None, ...], self.step)

def save_opts(self):
    """Save options to disk so we know what we ran this experiment with
    """
    models_dir = os.path.join(self.log_path, "models")
    if not os.path.exists(models_dir):
        os.makedirs(models_dir)
    to_save = self.opt.__dict__.copy()

    with open(os.path.join(models_dir, 'opt.json'), 'w') as f:
        json.dump(to_save, f, indent=2)

def save_model(self):
    """Save model weights to disk
    """
    save_folder = os.path.join(self.log_path, "models", "weights_{}".format(self.epoch))
    if not os.path.exists(save_folder):
        os.makedirs(save_folder)

    for model_name, model in self.models.items():
        save_path = os.path.join(save_folder, "{}.pth".format(model_name))
        to_save = model.state_dict()
        if model_name == 'encoder':
            # save the sizes - these are needed at prediction time
            to_save['height'] = self.opt.height
            to_save['width'] = self.opt.width
            to_save['use_stereo'] = self.opt.use_stereo
        torch.save(to_save, save_path)

    for model_name, model in self.models_pose.items():
        save_path = os.path.join(save_folder, "{}.pth".format(model_name))
        to_save = model.state_dict()
        torch.save(to_save, save_path)

    save_path = os.path.join(save_folder, "{}.pth".format("adam"))
    torch.save(self.model_optimizer.state_dict(), save_path)

    save_path = os.path.join(save_folder, "{}.pth".format("adam_pose"))
    if self.use_pose_net:
        torch.save(self.model_pose_optimizer.state_dict(), save_path)

def load_pretrain(self):
    self.opt.mypretrain = os.path.expanduser(self.opt.mypretrain)
    path = self.opt.mypretrain
    model_dict = self.models["encoder"].state_dict()
    pretrained_dict = torch.load(path)['model']
    pretrained_dict = {k: v for k, v in pretrained_dict.items() if (k in model_dict and not k.startswith('encoder_'))}
    model_dict.update(pretrained_dict)
    self.models["encoder"].load_state_dict(model_dict)
    print('mypretrain loaded.')

def load_model(self):
    """Load model(s) from disk
    """
    self.opt.load_weights_folder = os.path.expanduser(self.opt.load_weights_folder)

    assert os.path.isdir(self.opt.load_weights_folder), \
        "Cannot find folder {}".format(self.opt.load_weights_folder)
    print("loading model from folder {}".format(self.opt.load_weights_folder))

    for n in self.opt.models_to_load:
        print("Loading {} weights...".format(n))
        path = os.path.join(self.opt.load_weights_folder, "{}.pth".format(n))

        if n in ['pose_encoder', 'pose']:
            model_dict = self.models_pose[n].state_dict()
            pretrained_dict = torch.load(path)
            pretrained_dict = {k: v for k, v in pretrained_dict.items() if k in model_dict}
            model_dict.update(pretrained_dict)
            self.models_pose[n].load_state_dict(model_dict)
        else:
            model_dict = self.models[n].state_dict()
            pretrained_dict = torch.load(path)
            pretrained_dict = {k: v for k, v in pretrained_dict.items() if k in model_dict}
            model_dict.update(pretrained_dict)
            self.models[n].load_state_dict(model_dict)

```

```

# loading adam state

optimizer_load_path = os.path.join(self.opt.load_weights_folder, "adam.pth")
optimizer_pose_load_path = os.path.join(self.opt.load_weights_folder, "adam_pose.pth")
if os.path.isfile(optimizer_load_path):
    print("Loading Adam weights")
    optimizer_dict = torch.load(optimizer_load_path)
    optimizer_pose_dict = torch.load(optimizer_pose_load_path)
    self.model_optimizer.load_state_dict(optimizer_dict)
    self.model_pose_optimizer.load_state_dict(optimizer_pose_dict)
else:
    print("Cannot find Adam weights so Adam is randomly initialized")

options = LiteMonoOptions()
opts = options.parse()
par_dir = os.getcwd()
print(par_dir)

trainer = Trainer(opts)
trainer.train()

```

/mnt/workspace/sunqiao/mymono

/home/pai/lib/python3.9/site-packages/torchvision/models/_utils.py:252: UserWarning: Accessing the model URLs via the internal dictionary of the module is deprecated since 0.13 and will be removed in 0.15. Please access them via the appropriate Weights Enum instead.

warnings.warn(

Training model named:

lite-mono

Models and tensorboard events files are saved to:

./tmp

Training is using:

cuda

Using split:

eigen_zhou

There are 45619 training items and 4562 validation items

Training

/mnt/workspace/sunqiao/mymono/networks/depth_encoder.py:35: UserWarning: __floordiv__ is deprecated, and its behavior will change in a future version of pytorch. It currently rounds toward 0 (like the 'trunc' function NOT 'floor'). This results in incorrect rounding for negative values. To keep the current behavior, use torch.div(a, b, rounding_mode='trunc'), or for actual floor division, use torch.div(a, b, rounding_mode='floor').

dim_t = self.temperature ** (2 * (dim_t // 2) / self.hidden_dim)

epoch	0		lr	0.000100		lr_p	0.000100		batch	0		examples/s:	1.1		loss:	0.20615		time e
lapsed: 00h00m04s time left: 00h00m00s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	250		examples/s:	104.3		loss:	0.12414		time e
lapsed: 00h00m13s time left: 07h41m26s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	500		examples/s:	103.4		loss:	0.23405		time e
lapsed: 00h00m23s time left: 06h34m19s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	750		examples/s:	100.9		loss:	0.62217		time e
lapsed: 00h00m33s time left: 06h11m04s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	1000		examples/s:	106.4		loss:	0.27663		time e
lapsed: 00h00m43s time left: 05h58m49s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	1250		examples/s:	99.6		loss:	0.23921		time e
lapsed: 00h00m53s time left: 05h51m50s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	1500		examples/s:	105.8		loss:	0.30374		time e
lapsed: 00h01m03s time left: 05h47m40s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	1750		examples/s:	91.6		loss:	0.32157		time e
lapsed: 00h01m12s time left: 05h44m05s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	2000		examples/s:	104.3		loss:	0.51153		time e
lapsed: 00h01m22s time left: 05h41m27s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	2250		examples/s:	103.6		loss:	0.34248		time e
lapsed: 00h01m32s time left: 05h39m11s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	2500		examples/s:	106.5		loss:	0.45199		time e
lapsed: 00h01m42s time left: 05h37m25s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	2750		examples/s:	102.5		loss:	0.36953		time e
lapsed: 00h01m51s time left: 05h35m43s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	3000		examples/s:	103.9		loss:	0.45303		time e
lapsed: 00h02m01s time left: 05h34m25s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	3250		examples/s:	101.6		loss:	0.50665		time e
lapsed: 00h02m11s time left: 05h33m13s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	3500		examples/s:	90.4		loss:	0.61825		time e
lapsed: 00h02m21s time left: 05h32m04s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	3750		examples/s:	106.5		loss:	0.60815		time e
lapsed: 00h02m30s time left: 05h31m12s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	4000		examples/s:	105.6		loss:	0.46399		time e
lapsed: 00h02m40s time left: 05h30m48s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	4250		examples/s:	104.2		loss:	0.75837		time e
lapsed: 00h02m50s time left: 05h29m48s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	4500		examples/s:	105.3		loss:	0.80678		time e
lapsed: 00h03m00s time left: 05h29m01s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	4750		examples/s:	104.7		loss:	0.39219		time e
lapsed: 00h03m09s time left: 05h28m30s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	5000		examples/s:	105.9		loss:	0.69590		time e
lapsed: 00h03m20s time left: 05h28m26s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	5250		examples/s:	101.2		loss:	0.23884		time e
lapsed: 00h03m29s time left: 05h27m51s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	5500		examples/s:	98.9		loss:	0.52098		time e
lapsed: 00h03m39s time left: 05h27m34s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	5750		examples/s:	101.5		loss:	0.37412		time e
lapsed: 00h03m49s time left: 05h27m06s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	6000		examples/s:	103.0		loss:	0.97765		time e
lapsed: 00h03m59s time left: 05h26m38s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	6250		examples/s:	104.0		loss:	0.22806		time e
lapsed: 00h04m09s time left: 05h26m15s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	6500		examples/s:	100.6		loss:	0.43808		time e
lapsed: 00h04m18s time left: 05h25m50s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	6750		examples/s:	100.8		loss:	0.54961		time e
lapsed: 00h04m28s time left: 05h25m25s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	7000		examples/s:	104.1		loss:	0.53456		time e
lapsed: 00h04m38s time left: 05h25m03s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	7250		examples/s:	104.8		loss:	0.47826		time e
lapsed: 00h04m48s time left: 05h24m44s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	7500		examples/s:	103.3		loss:	0.35346		time e
lapsed: 00h04m57s time left: 05h24m26s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	7750		examples/s:	104.4		loss:	0.14338		time e
lapsed: 00h05m07s time left: 05h24m14s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	8000		examples/s:	101.7		loss:	0.38821		time e
lapsed: 00h05m17s time left: 05h23m56s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	8250		examples/s:	90.9		loss:	0.37182		time e
lapsed: 00h05m27s time left: 05h23m36s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	8500		examples/s:	104.3		loss:	0.39300		time e
lapsed: 00h05m37s time left: 05h23m16s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	8750		examples/s:	106.5		loss:	0.54625		time e
lapsed: 00h05m46s time left: 05h22m54s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	9000		examples/s:	101.1		loss:	0.68417		time e
lapsed: 00h05m56s time left: 05h22m36s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	9250		examples/s:	104.8		loss:	0.50373		time e
lapsed: 00h06m06s time left: 05h22m20s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	9500		examples/s:	104.6		loss:	0.48334		time e
lapsed: 00h06m16s time left: 05h22m06s																		
epoch	0		lr	0.000100		lr_p	0.000100		batch	9750		examples/s:	106.1		loss:	0.59760		time e
lapsed: 00h06m25s time left: 05h21m52s																		

epoch	0		lr 0.000100		lr_p 0.000100		batch	10000		examples/s: 103.0		loss: 0.31928		time e
lapsed: 00h06m35s time left: 05h21m37s														
epoch	0		lr 0.000100		lr_p 0.000100		batch	10250		examples/s: 99.1		loss: 0.65071		time e
lapsed: 00h06m45s time left: 05h21m26s														
epoch	0		lr 0.000100		lr_p 0.000100		batch	10500		examples/s: 108.1		loss: 0.72714		time e
lapsed: 00h06m55s time left: 05h21m14s														
epoch	0		lr 0.000100		lr_p 0.000100		batch	10750		examples/s: 100.5		loss: 0.39611		time e
lapsed: 00h07m05s time left: 05h20m57s														
epoch	0		lr 0.000100		lr_p 0.000100		batch	11000		examples/s: 104.1		loss: 0.52416		time e
lapsed: 00h07m15s time left: 05h20m53s														
epoch	0		lr 0.000100		lr_p 0.000100		batch	11250		examples/s: 102.8		loss: 0.41693		time e
lapsed: 00h07m25s time left: 05h20m40s														
Training														
epoch	1		lr 0.000099		lr_p 0.000099		batch	0		examples/s: 21.1		loss: 0.55769		time e
lapsed: 00h07m33s time left: 05h22m05s														
epoch	1		lr 0.000099		lr_p 0.000099		batch	250		examples/s: 94.2		loss: 0.38379		time e
lapsed: 00h07m43s time left: 05h21m57s														
epoch	1		lr 0.000099		lr_p 0.000099		batch	500		examples/s: 101.7		loss: 0.57161		time e
lapsed: 00h07m53s time left: 05h21m42s														
epoch	1		lr 0.000099		lr_p 0.000099		batch	596		examples/s: 99.8		loss: 0.87279		time e
lapsed: 00h07m56s time left: 05h21m38s														
epoch	1		lr 0.000099		lr_p 0.000099		batch	750		examples/s: 102.4		loss: 0.41252		time e
lapsed: 00h08m02s time left: 05h21m30s														
epoch	1		lr 0.000099		lr_p 0.000099		batch	1000		examples/s: 98.9		loss: 0.39082		time e
lapsed: 00h08m12s time left: 05h21m16s														
epoch	1		lr 0.000099		lr_p 0.000099		batch	1250		examples/s: 99.1		loss: 0.38651		time e
lapsed: 00h08m22s time left: 05h21m03s														
epoch	1		lr 0.000099		lr_p 0.000099		batch	1500		examples/s: 104.1		loss: 0.46981		time e
lapsed: 00h08m32s time left: 05h20m51s														
epoch	1		lr 0.000099		lr_p 0.000099		batch	1750		examples/s: 107.3		loss: 0.70182		time e
lapsed: 00h08m42s time left: 05h20m40s														
epoch	1		lr 0.000099		lr_p 0.000099		batch	2000		examples/s: 105.0		loss: 0.52628		time e
lapsed: 00h08m52s time left: 05h20m30s														
epoch	1		lr 0.000099		lr_p 0.000099		batch	2250		examples/s: 106.0		loss: 0.54541		time e
lapsed: 00h09m02s time left: 05h20m17s														
epoch	1		lr 0.000099		lr_p 0.000099		batch	2500		examples/s: 99.9		loss: 0.33003		time e
lapsed: 00h09m12s time left: 05h20m07s														
epoch	1		lr 0.000099		lr_p 0.000099		batch	2596		examples/s: 101.9		loss: 0.60328		time e
lapsed: 00h09m15s time left: 05h20m03s														
epoch	1		lr 0.000099		lr_p 0.000099		batch	2750		examples/s: 100.7		loss: 0.96149		time e
lapsed: 00h09m22s time left: 05h19m57s														
epoch	1		lr 0.000099		lr_p 0.000099		batch	3000		examples/s: 107.0		loss: 0.67061		time e
lapsed: 00h09m32s time left: 05h19m52s														
epoch	1		lr 0.000099		lr_p 0.000099		batch	3250		examples/s: 102.5		loss: 0.55636		time e
lapsed: 00h09m41s time left: 05h19m39s														
epoch	1		lr 0.000099		lr_p 0.000099		batch	3500		examples/s: 93.7		loss: 0.56087		time e
lapsed: 00h09m51s time left: 05h19m24s														
epoch	1		lr 0.000099		lr_p 0.000099		batch	3750		examples/s: 106.1		loss: 0.24713		time e
lapsed: 00h10m01s time left: 05h19m12s														
epoch	1		lr 0.000099		lr_p 0.000099		batch	4000		examples/s: 105.0		loss: 0.28251		time e
lapsed: 00h10m11s time left: 05h19m00s														
epoch	1		lr 0.000099		lr_p 0.000099		batch	4250		examples/s: 98.6		loss: 0.82811		time e
lapsed: 00h10m21s time left: 05h18m51s														
epoch	1		lr 0.000099		lr_p 0.000099		batch	4500		examples/s: 105.8		loss: 0.80209		time e
lapsed: 00h10m31s time left: 05h18m41s														
epoch	1		lr 0.000099		lr_p 0.000099		batch	4596		examples/s: 103.1		loss: 1.03606		time e
lapsed: 00h10m35s time left: 05h18m35s														
epoch	1		lr 0.000099		lr_p 0.000099		batch	4750		examples/s: 104.1		loss: 0.48541		time e
lapsed: 00h10m41s time left: 05h18m29s														
epoch	1		lr 0.000099		lr_p 0.000099		batch	5000		examples/s: 92.9		loss: 0.43020		time e
lapsed: 00h10m51s time left: 05h18m17s														
epoch	1		lr 0.000099		lr_p 0.000099		batch	5250		examples/s: 99.8		loss: 0.63685		time e
lapsed: 00h11m00s time left: 05h18m04s														
epoch	1		lr 0.000099		lr_p 0.000099		batch	5500		examples/s: 96.7		loss: 0.92930		time e
lapsed: 00h11m10s time left: 05h17m51s														
epoch	1		lr 0.000099		lr_p 0.000099		batch	5750		examples/s: 102.4		loss: 0.59154		time e
lapsed: 00h11m20s time left: 05h17m39s														
epoch	1		lr 0.000099		lr_p 0.000099		batch	6000		examples/s: 103.0		loss: 0.54017		time e
lapsed: 00h11m30s time left: 05h17m27s														
epoch	1		lr 0.000099		lr_p 0.000099		batch	6250		examples/s: 104.0		loss: 0.45813		time e
lapsed: 00h11m40s time left: 05h17m16s														
epoch	1		lr 0.000099		lr_p 0.000099		batch	6500		examples/s: 102.6		loss: 0.52609		time e
lapsed: 00h11m50s time left: 05h17m05s														
epoch	1		lr 0.000099		lr_p 0.000099		batch	6596		examples/s: 103.4		loss: 0.48732		time e
lapsed: 00h11m53s time left: 05h17m00s														
epoch	1		lr 0.000099		lr_p 0.000099		batch	6750		examples/s: 97.6		loss: 0.43162		time e
lapsed: 00h11m59s time left: 05h16m53s														
epoch	1		lr 0.000099		lr_p 0.000099		batch	7000		examples/s: 98.1		loss: 0.60409		time e
lapsed: 00h12m09s time left: 05h16m42s														
epoch	1		lr 0.000099		lr_p 0.000099		batch	7250		examples/s: 102.1		loss: 0.50888		time e

```
lapsed: 00h12m19s | time left: 05h16m32s
epoch 1 | lr 0.000099 |lr_p 0.000099 | batch 7500 | examples/s: 101.5 | loss: 0.33882 | time e
lapsed: 00h12m29s | time left: 05h16m21s
epoch 1 | lr 0.000099 |lr_p 0.000099 | batch 7750 | examples/s: 102.2 | loss: 0.82656 | time e
lapsed: 00h12m39s | time left: 05h16m09s
epoch 1 | lr 0.000099 |lr_p 0.000099 | batch 8000 | examples/s: 99.8 | loss: 0.33599 | time e
lapsed: 00h12m49s | time left: 05h15m59s
epoch 1 | lr 0.000099 |lr_p 0.000099 | batch 8250 | examples/s: 103.2 | loss: 0.45911 | time e
lapsed: 00h12m59s | time left: 05h15m49s
epoch 1 | lr 0.000099 |lr_p 0.000099 | batch 8500 | examples/s: 102.0 | loss: 0.71029 | time e
lapsed: 00h13m09s | time left: 05h15m38s
epoch 1 | lr 0.000099 |lr_p 0.000099 | batch 8596 | examples/s: 94.3 | loss: 0.51233 | time e
lapsed: 00h13m12s | time left: 05h15m34s
epoch 1 | lr 0.000099 |lr_p 0.000099 | batch 10596 | examples/s: 105.9 | loss: 0.33377 | time e
lapsed: 00h14m31s | time left: 05h14m03s
Training
epoch 2 | lr 0.000098 |lr_p 0.000098 | batch 1192 | examples/s: 101.1 | loss: 0.34211 | time e
lapsed: 00h15m52s | time left: 05h13m17s
epoch 2 | lr 0.000098 |lr_p 0.000098 | batch 3192 | examples/s: 99.2 | loss: 0.85638 | time e
lapsed: 00h17m11s | time left: 05h11m46s
epoch 2 | lr 0.000098 |lr_p 0.000098 | batch 5192 | examples/s: 103.6 | loss: 0.31263 | time e
lapsed: 00h18m29s | time left: 05h10m10s
epoch 2 | lr 0.000098 |lr_p 0.000098 | batch 7192 | examples/s: 102.8 | loss: 0.57371 | time e
lapsed: 00h19m47s | time left: 05h08m26s
epoch 2 | lr 0.000098 |lr_p 0.000098 | batch 9192 | examples/s: 107.2 | loss: 0.98736 | time e
lapsed: 00h21m05s | time left: 05h06m50s
epoch 2 | lr 0.000098 |lr_p 0.000098 | batch 11192 | examples/s: 96.6 | loss: 0.44031 | time e
lapsed: 00h22m23s | time left: 05h05m15s
Training
epoch 3 | lr 0.000096 |lr_p 0.000096 | batch 1788 | examples/s: 99.7 | loss: 0.51225 | time e
lapsed: 00h23m45s | time left: 05h04m32s
epoch 3 | lr 0.000096 |lr_p 0.000096 | batch 3788 | examples/s: 100.5 | loss: 0.34387 | time e
lapsed: 00h25m03s | time left: 05h03m09s
epoch 3 | lr 0.000096 |lr_p 0.000096 | batch 5788 | examples/s: 104.8 | loss: 0.41121 | time e
lapsed: 00h26m22s | time left: 05h01m44s
epoch 3 | lr 0.000096 |lr_p 0.000096 | batch 7788 | examples/s: 102.5 | loss: 0.64221 | time e
lapsed: 00h27m40s | time left: 05h00m17s
epoch 3 | lr 0.000096 |lr_p 0.000096 | batch 9788 | examples/s: 99.8 | loss: 0.56477 | time e
lapsed: 00h29m00s | time left: 04h59m05s
Training
epoch 4 | lr 0.000094 |lr_p 0.000094 | batch 384 | examples/s: 99.7 | loss: 0.36104 | time e
lapsed: 00h30m22s | time left: 04h58m11s
epoch 4 | lr 0.000094 |lr_p 0.000094 | batch 2384 | examples/s: 104.5 | loss: 0.62724 | time e
lapsed: 00h31m41s | time left: 04h56m50s
epoch 4 | lr 0.000094 |lr_p 0.000094 | batch 4384 | examples/s: 103.9 | loss: 0.44416 | time e
lapsed: 00h33m00s | time left: 04h55m27s
epoch 4 | lr 0.000094 |lr_p 0.000094 | batch 6384 | examples/s: 101.2 | loss: 0.57530 | time e
lapsed: 00h34m19s | time left: 04h54m07s
epoch 4 | lr 0.000094 |lr_p 0.000094 | batch 8384 | examples/s: 103.8 | loss: 0.32936 | time e
lapsed: 00h35m39s | time left: 04h52m51s
epoch 4 | lr 0.000094 |lr_p 0.000094 | batch 10384 | examples/s: 94.6 | loss: 0.52077 | time e
lapsed: 00h36m57s | time left: 04h51m29s
Training
epoch 5 | lr 0.000091 |lr_p 0.000092 | batch 980 | examples/s: 96.4 | loss: 0.56453 | time e
lapsed: 00h38m19s | time left: 04h50m29s
epoch 5 | lr 0.000091 |lr_p 0.000092 | batch 2980 | examples/s: 102.1 | loss: 0.51823 | time e
lapsed: 00h39m38s | time left: 04h49m10s
epoch 5 | lr 0.000091 |lr_p 0.000092 | batch 4980 | examples/s: 102.4 | loss: 0.43740 | time e
lapsed: 00h40m57s | time left: 04h47m48s
epoch 5 | lr 0.000091 |lr_p 0.000092 | batch 6980 | examples/s: 103.3 | loss: 0.65279 | time e
lapsed: 00h42m18s | time left: 04h46m34s
epoch 5 | lr 0.000091 |lr_p 0.000092 | batch 8980 | examples/s: 95.9 | loss: 0.76437 | time e
lapsed: 00h43m37s | time left: 04h45m18s
epoch 5 | lr 0.000091 |lr_p 0.000092 | batch 10980 | examples/s: 97.3 | loss: 0.37074 | time e
lapsed: 00h44m57s | time left: 04h44m03s
Training
epoch 6 | lr 0.000089 |lr_p 0.000089 | batch 1576 | examples/s: 95.4 | loss: 0.42072 | time e
lapsed: 00h46m19s | time left: 04h42m58s
epoch 6 | lr 0.000089 |lr_p 0.000089 | batch 3576 | examples/s: 106.7 | loss: 0.57686 | time e
lapsed: 00h47m39s | time left: 04h41m41s
epoch 6 | lr 0.000089 |lr_p 0.000089 | batch 5576 | examples/s: 90.7 | loss: 0.37693 | time e
lapsed: 00h48m58s | time left: 04h40m20s
epoch 6 | lr 0.000089 |lr_p 0.000089 | batch 7576 | examples/s: 96.7 | loss: 0.47295 | time e
lapsed: 00h50m18s | time left: 04h39m05s
epoch 6 | lr 0.000089 |lr_p 0.000089 | batch 9576 | examples/s: 100.4 | loss: 0.27683 | time e
lapsed: 00h51m38s | time left: 04h37m46s
Training
epoch 7 | lr 0.000085 |lr_p 0.000086 | batch 172 | examples/s: 105.0 | loss: 0.32946 | time e
lapsed: 00h53m00s | time left: 04h36m40s
epoch 7 | lr 0.000085 |lr_p 0.000086 | batch 2172 | examples/s: 100.1 | loss: 0.65826 | time e
```

```
lapsed: 00h54m19s | time left: 04h35m18s
epoch 7 | lr 0.000085 | lr_p 0.000086 | batch 4172 | examples/s: 99.8 | loss: 0.23194 | time e
lapsed: 00h55m38s | time left: 04h33m59s
epoch 7 | lr 0.000085 | lr_p 0.000086 | batch 6172 | examples/s: 95.7 | loss: 0.61961 | time e
lapsed: 00h56m58s | time left: 04h32m41s
epoch 7 | lr 0.000085 | lr_p 0.000086 | batch 8172 | examples/s: 101.6 | loss: 0.38886 | time e
lapsed: 00h58m18s | time left: 04h31m22s
epoch 7 | lr 0.000085 | lr_p 0.000086 | batch 10172 | examples/s: 102.1 | loss: 0.39918 | time e
lapsed: 00h59m37s | time left: 04h30m02s
Training
epoch 8 | lr 0.000082 | lr_p 0.000083 | batch 768 | examples/s: 103.7 | loss: 0.37604 | time e
lapsed: 01h00m58s | time left: 04h28m50s
epoch 8 | lr 0.000082 | lr_p 0.000083 | batch 2768 | examples/s: 96.8 | loss: 0.33777 | time e
lapsed: 01h02m19s | time left: 04h27m36s
epoch 8 | lr 0.000082 | lr_p 0.000083 | batch 4768 | examples/s: 99.1 | loss: 0.26284 | time e
lapsed: 01h03m39s | time left: 04h26m16s
epoch 8 | lr 0.000082 | lr_p 0.000083 | batch 6768 | examples/s: 104.3 | loss: 0.71091 | time e
lapsed: 01h04m59s | time left: 04h24m58s
epoch 8 | lr 0.000082 | lr_p 0.000083 | batch 8768 | examples/s: 104.4 | loss: 0.37017 | time e
lapsed: 01h06m18s | time left: 04h23m40s
epoch 8 | lr 0.000082 | lr_p 0.000083 | batch 10768 | examples/s: 96.0 | loss: 0.28578 | time e
lapsed: 01h07m37s | time left: 04h22m17s
Training
epoch 9 | lr 0.000078 | lr_p 0.000079 | batch 1364 | examples/s: 104.9 | loss: 0.30243 | time e
lapsed: 01h08m58s | time left: 04h21m04s
epoch 9 | lr 0.000078 | lr_p 0.000079 | batch 3364 | examples/s: 101.7 | loss: 0.36172 | time e
lapsed: 01h10m18s | time left: 04h19m44s
epoch 9 | lr 0.000078 | lr_p 0.000079 | batch 5364 | examples/s: 104.3 | loss: 0.42035 | time e
lapsed: 01h11m33s | time left: 04h18m08s
epoch 9 | lr 0.000078 | lr_p 0.000079 | batch 7364 | examples/s: 112.3 | loss: 0.46136 | time e
lapsed: 01h12m44s | time left: 04h16m17s
epoch 9 | lr 0.000078 | lr_p 0.000079 | batch 9364 | examples/s: 115.1 | loss: 1.12097 | time e
lapsed: 01h13m54s | time left: 04h14m27s
epoch 9 | lr 0.000078 | lr_p 0.000079 | batch 11364 | examples/s: 109.6 | loss: 0.50139 | time e
lapsed: 01h15m05s | time left: 04h12m39s
Training
epoch 10 | lr 0.000073 | lr_p 0.000075 | batch 1960 | examples/s: 111.4 | loss: 0.52033 | time e
lapsed: 01h16m18s | time left: 04h11m01s
epoch 10 | lr 0.000073 | lr_p 0.000075 | batch 3960 | examples/s: 113.3 | loss: 0.73017 | time e
lapsed: 01h17m29s | time left: 04h09m18s
epoch 10 | lr 0.000073 | lr_p 0.000075 | batch 5960 | examples/s: 116.0 | loss: 0.39833 | time e
lapsed: 01h18m41s | time left: 04h07m36s
epoch 10 | lr 0.000073 | lr_p 0.000075 | batch 7960 | examples/s: 110.1 | loss: 0.46344 | time e
lapsed: 01h19m53s | time left: 04h05m57s
epoch 10 | lr 0.000073 | lr_p 0.000075 | batch 9960 | examples/s: 117.6 | loss: 0.38397 | time e
lapsed: 01h21m05s | time left: 04h04m18s
Training
epoch 11 | lr 0.000069 | lr_p 0.000071 | batch 556 | examples/s: 111.0 | loss: 0.59284 | time e
lapsed: 01h22m19s | time left: 04h02m46s
epoch 11 | lr 0.000069 | lr_p 0.000071 | batch 2556 | examples/s: 111.8 | loss: 0.40175 | time e
lapsed: 01h23m32s | time left: 04h01m13s
epoch 11 | lr 0.000069 | lr_p 0.000071 | batch 4556 | examples/s: 112.1 | loss: 0.40041 | time e
lapsed: 01h24m45s | time left: 03h59m41s
epoch 11 | lr 0.000069 | lr_p 0.000071 | batch 6556 | examples/s: 108.3 | loss: 0.38449 | time e
lapsed: 01h25m58s | time left: 03h58m08s
epoch 11 | lr 0.000069 | lr_p 0.000071 | batch 8556 | examples/s: 103.2 | loss: 0.51516 | time e
lapsed: 01h27m11s | time left: 03h56m34s
epoch 11 | lr 0.000069 | lr_p 0.000071 | batch 10556 | examples/s: 111.2 | loss: 0.76401 | time e
lapsed: 01h28m23s | time left: 03h55m00s
Training
epoch 12 | lr 0.000064 | lr_p 0.000066 | batch 1152 | examples/s: 106.6 | loss: 0.94782 | time e
lapsed: 01h29m38s | time left: 03h53m34s
epoch 12 | lr 0.000064 | lr_p 0.000066 | batch 3152 | examples/s: 112.9 | loss: 0.52948 | time e
lapsed: 01h30m50s | time left: 03h52m02s
epoch 12 | lr 0.000064 | lr_p 0.000066 | batch 5152 | examples/s: 92.4 | loss: 0.46954 | time e
lapsed: 01h32m03s | time left: 03h50m31s
epoch 12 | lr 0.000064 | lr_p 0.000066 | batch 7152 | examples/s: 114.0 | loss: 0.38311 | time e
lapsed: 01h33m15s | time left: 03h48m59s
epoch 12 | lr 0.000064 | lr_p 0.000066 | batch 9152 | examples/s: 115.5 | loss: 0.45565 | time e
lapsed: 01h34m27s | time left: 03h47m27s
epoch 12 | lr 0.000064 | lr_p 0.000066 | batch 11152 | examples/s: 116.9 | loss: 0.37749 | time e
lapsed: 01h35m38s | time left: 03h45m56s
Training
epoch 13 | lr 0.000060 | lr_p 0.000062 | batch 1748 | examples/s: 113.2 | loss: 0.46985 | time e
lapsed: 01h36m53s | time left: 03h44m31s
epoch 13 | lr 0.000060 | lr_p 0.000062 | batch 3748 | examples/s: 114.1 | loss: 0.54206 | time e
lapsed: 01h38m06s | time left: 03h43m04s
epoch 13 | lr 0.000060 | lr_p 0.000062 | batch 5748 | examples/s: 109.0 | loss: 0.22653 | time e
lapsed: 01h39m19s | time left: 03h41m35s
epoch 13 | lr 0.000060 | lr_p 0.000062 | batch 7748 | examples/s: 115.8 | loss: 0.65166 | time e
```

```

lapsed: 01h40m30s | time left: 03h40m06s
epoch 13 | lr 0.000060 |lr_p 0.000062 | batch 9748 | examples/s: 112.4 | loss: 0.43550 | time e
lapsed: 01h41m42s | time left: 03h38m36s
Training
epoch 14 | lr 0.000055 |lr_p 0.000057 | batch 344 | examples/s: 106.5 | loss: 0.35515 | time e
lapsed: 01h42m57s | time left: 03h37m13s
epoch 14 | lr 0.000055 |lr_p 0.000057 | batch 2344 | examples/s: 109.8 | loss: 0.22190 | time e
lapsed: 01h44m11s | time left: 03h35m49s
epoch 14 | lr 0.000055 |lr_p 0.000057 | batch 4344 | examples/s: 106.4 | loss: 0.55782 | time e
lapsed: 01h45m23s | time left: 03h34m23s
epoch 14 | lr 0.000055 |lr_p 0.000057 | batch 6344 | examples/s: 111.4 | loss: 0.32657 | time e
lapsed: 01h46m35s | time left: 03h32m56s
epoch 14 | lr 0.000055 |lr_p 0.000057 | batch 8344 | examples/s: 113.8 | loss: 0.40553 | time e
lapsed: 01h47m47s | time left: 03h31m28s
epoch 14 | lr 0.000055 |lr_p 0.000057 | batch 10344 | examples/s: 112.1 | loss: 0.51742 | time e
lapsed: 01h48m58s | time left: 03h30m00s
Training
epoch 15 | lr 0.000050 |lr_p 0.000053 | batch 940 | examples/s: 105.7 | loss: 0.56769 | time e
lapsed: 01h50m13s | time left: 03h28m40s
epoch 15 | lr 0.000050 |lr_p 0.000053 | batch 2940 | examples/s: 103.5 | loss: 0.34662 | time e
lapsed: 01h51m26s | time left: 03h27m15s
epoch 15 | lr 0.000050 |lr_p 0.000053 | batch 4940 | examples/s: 92.9 | loss: 0.53990 | time e
lapsed: 01h52m39s | time left: 03h25m50s
epoch 15 | lr 0.000050 |lr_p 0.000053 | batch 6940 | examples/s: 113.8 | loss: 0.43283 | time e
lapsed: 01h53m50s | time left: 03h24m24s
epoch 15 | lr 0.000050 |lr_p 0.000053 | batch 8940 | examples/s: 113.4 | loss: 0.39722 | time e
lapsed: 01h55m03s | time left: 03h23m00s
epoch 15 | lr 0.000050 |lr_p 0.000053 | batch 10940 | examples/s: 111.0 | loss: 0.66041 | time e
lapsed: 01h56m15s | time left: 03h21m35s
Training
epoch 16 | lr 0.000045 |lr_p 0.000048 | batch 1536 | examples/s: 105.4 | loss: 0.56372 | time e
lapsed: 01h57m30s | time left: 03h20m15s
epoch 16 | lr 0.000045 |lr_p 0.000048 | batch 3536 | examples/s: 113.1 | loss: 0.31197 | time e
lapsed: 01h58m43s | time left: 03h18m53s
epoch 16 | lr 0.000045 |lr_p 0.000048 | batch 5536 | examples/s: 107.3 | loss: 0.47659 | time e
lapsed: 01h59m55s | time left: 03h17m30s
epoch 16 | lr 0.000045 |lr_p 0.000048 | batch 7536 | examples/s: 112.9 | loss: 0.63046 | time e
lapsed: 02h01m09s | time left: 03h16m08s
epoch 16 | lr 0.000045 |lr_p 0.000048 | batch 9536 | examples/s: 111.3 | loss: 0.57245 | time e
lapsed: 02h02m21s | time left: 03h14m44s
Training
epoch 17 | lr 0.000041 |lr_p 0.000044 | batch 132 | examples/s: 110.6 | loss: 0.50303 | time e
lapsed: 02h03m34s | time left: 03h13m24s
epoch 17 | lr 0.000041 |lr_p 0.000044 | batch 2132 | examples/s: 106.9 | loss: 0.39004 | time e
lapsed: 02h04m48s | time left: 03h12m03s
epoch 17 | lr 0.000041 |lr_p 0.000044 | batch 4132 | examples/s: 105.1 | loss: 0.71872 | time e
lapsed: 02h06m01s | time left: 03h10m41s
epoch 17 | lr 0.000041 |lr_p 0.000044 | batch 6132 | examples/s: 114.0 | loss: 0.67266 | time e
lapsed: 02h07m15s | time left: 03h09m21s
epoch 17 | lr 0.000041 |lr_p 0.000044 | batch 8132 | examples/s: 113.8 | loss: 0.37485 | time e
lapsed: 02h08m28s | time left: 03h08m00s
epoch 17 | lr 0.000041 |lr_p 0.000044 | batch 10132 | examples/s: 96.1 | loss: 0.23352 | time e
lapsed: 02h09m40s | time left: 03h06m38s
Training
epoch 18 | lr 0.000036 |lr_p 0.000039 | batch 728 | examples/s: 108.7 | loss: 0.30353 | time e
lapsed: 02h10m55s | time left: 03h05m19s
epoch 18 | lr 0.000036 |lr_p 0.000039 | batch 2728 | examples/s: 111.0 | loss: 0.74875 | time e
lapsed: 02h12m08s | time left: 03h03m58s
epoch 18 | lr 0.000036 |lr_p 0.000039 | batch 4728 | examples/s: 109.7 | loss: 0.55153 | time e
lapsed: 02h13m21s | time left: 03h02m37s
epoch 18 | lr 0.000036 |lr_p 0.000039 | batch 6728 | examples/s: 91.7 | loss: 0.45575 | time e
lapsed: 02h14m33s | time left: 03h01m16s
epoch 18 | lr 0.000036 |lr_p 0.000039 | batch 8728 | examples/s: 107.3 | loss: 0.32482 | time e
lapsed: 02h15m45s | time left: 02h59m54s
epoch 18 | lr 0.000036 |lr_p 0.000039 | batch 10728 | examples/s: 113.2 | loss: 0.29835 | time e
lapsed: 02h16m57s | time left: 02h58m33s
Training
epoch 19 | lr 0.000032 |lr_p 0.000035 | batch 1324 | examples/s: 114.9 | loss: 0.44576 | time e
lapsed: 02h18m12s | time left: 02h57m16s
epoch 19 | lr 0.000032 |lr_p 0.000035 | batch 3324 | examples/s: 110.6 | loss: 0.53154 | time e
lapsed: 02h19m25s | time left: 02h55m55s
epoch 19 | lr 0.000032 |lr_p 0.000035 | batch 5324 | examples/s: 115.5 | loss: 0.78353 | time e
lapsed: 02h20m37s | time left: 02h54m34s
epoch 19 | lr 0.000032 |lr_p 0.000035 | batch 7324 | examples/s: 106.3 | loss: 0.72005 | time e
lapsed: 02h21m50s | time left: 02h53m14s
epoch 19 | lr 0.000032 |lr_p 0.000035 | batch 9324 | examples/s: 106.5 | loss: 0.29777 | time e
lapsed: 02h23m02s | time left: 02h51m53s
epoch 19 | lr 0.000032 |lr_p 0.000035 | batch 11324 | examples/s: 114.4 | loss: 0.31344 | time e
lapsed: 02h24m14s | time left: 02h50m33s
Training

```



```

epoch 20 | lr 0.000027 |lr_p 0.000031 | batch 1920 | examples/s: 114.1 | loss: 0.30517 | time e
lapsed: 02h25m29s | time left: 02h49m16s
epoch 20 | lr 0.000027 |lr_p 0.000031 | batch 3920 | examples/s: 112.8 | loss: 0.62174 | time e
lapsed: 02h26m42s | time left: 02h47m56s
epoch 20 | lr 0.000027 |lr_p 0.000031 | batch 5920 | examples/s: 114.6 | loss: 0.15017 | time e
lapsed: 02h27m54s | time left: 02h46m36s
epoch 20 | lr 0.000027 |lr_p 0.000031 | batch 7920 | examples/s: 112.7 | loss: 0.35059 | time e
lapsed: 02h29m06s | time left: 02h45m16s
epoch 20 | lr 0.000027 |lr_p 0.000031 | batch 9920 | examples/s: 112.9 | loss: 0.60489 | time e
lapsed: 02h30m17s | time left: 02h43m56s
Training
epoch 21 | lr 0.000023 |lr_p 0.000027 | batch 516 | examples/s: 114.3 | loss: 0.25326 | time e
lapsed: 02h31m32s | time left: 02h42m39s
epoch 21 | lr 0.000023 |lr_p 0.000027 | batch 2516 | examples/s: 114.3 | loss: 0.47149 | time e
lapsed: 02h32m44s | time left: 02h41m19s
epoch 21 | lr 0.000023 |lr_p 0.000027 | batch 4516 | examples/s: 109.7 | loss: 0.51162 | time e
lapsed: 02h33m56s | time left: 02h40m00s
epoch 21 | lr 0.000023 |lr_p 0.000027 | batch 6516 | examples/s: 114.7 | loss: 0.49498 | time e
lapsed: 02h35m09s | time left: 02h38m41s
epoch 21 | lr 0.000023 |lr_p 0.000027 | batch 8516 | examples/s: 112.8 | loss: 0.23164 | time e
lapsed: 02h36m21s | time left: 02h37m21s
epoch 21 | lr 0.000023 |lr_p 0.000027 | batch 10516 | examples/s: 113.7 | loss: 0.57831 | time e
lapsed: 02h37m33s | time left: 02h36m02s
Training
epoch 22 | lr 0.000020 |lr_p 0.000024 | batch 1112 | examples/s: 108.8 | loss: 0.71911 | time e
lapsed: 02h38m48s | time left: 02h34m46s
epoch 22 | lr 0.000020 |lr_p 0.000024 | batch 3112 | examples/s: 107.6 | loss: 0.46787 | time e
lapsed: 02h39m59s | time left: 02h33m26s
epoch 22 | lr 0.000020 |lr_p 0.000024 | batch 5112 | examples/s: 112.8 | loss: 0.45450 | time e
lapsed: 02h41m11s | time left: 02h32m07s
epoch 22 | lr 0.000020 |lr_p 0.000024 | batch 7112 | examples/s: 112.6 | loss: 0.62747 | time e
lapsed: 02h42m23s | time left: 02h30m48s
epoch 22 | lr 0.000020 |lr_p 0.000024 | batch 9112 | examples/s: 115.6 | loss: 0.57191 | time e
lapsed: 02h43m35s | time left: 02h29m30s
epoch 22 | lr 0.000020 |lr_p 0.000024 | batch 11112 | examples/s: 113.8 | loss: 0.42776 | time e
lapsed: 02h44m48s | time left: 02h28m11s
Training
epoch 23 | lr 0.000016 |lr_p 0.000021 | batch 1708 | examples/s: 106.1 | loss: 0.72683 | time e
lapsed: 02h46m03s | time left: 02h26m56s
epoch 23 | lr 0.000016 |lr_p 0.000021 | batch 3708 | examples/s: 116.4 | loss: 0.54990 | time e
lapsed: 02h47m15s | time left: 02h25m37s
epoch 23 | lr 0.000016 |lr_p 0.000021 | batch 5708 | examples/s: 115.7 | loss: 0.41523 | time e
lapsed: 02h48m26s | time left: 02h24m18s
epoch 23 | lr 0.000016 |lr_p 0.000021 | batch 7708 | examples/s: 113.6 | loss: 0.36754 | time e
lapsed: 02h49m38s | time left: 02h22m59s
epoch 23 | lr 0.000016 |lr_p 0.000021 | batch 9708 | examples/s: 111.9 | loss: 0.33181 | time e
lapsed: 02h50m50s | time left: 02h21m41s
Training
epoch 24 | lr 0.000014 |lr_p 0.000018 | batch 304 | examples/s: 112.1 | loss: 0.64990 | time e
lapsed: 02h52m04s | time left: 02h20m25s
epoch 24 | lr 0.000014 |lr_p 0.000018 | batch 2304 | examples/s: 109.3 | loss: 0.46897 | time e
lapsed: 02h53m18s | time left: 02h19m08s
epoch 24 | lr 0.000014 |lr_p 0.000018 | batch 4304 | examples/s: 111.0 | loss: 0.57997 | time e
lapsed: 02h54m31s | time left: 02h17m51s
epoch 24 | lr 0.000014 |lr_p 0.000018 | batch 6304 | examples/s: 109.0 | loss: 0.75668 | time e
lapsed: 02h55m43s | time left: 02h16m34s
epoch 24 | lr 0.000014 |lr_p 0.000018 | batch 8304 | examples/s: 111.1 | loss: 0.34847 | time e
lapsed: 02h56m56s | time left: 02h15m16s
epoch 24 | lr 0.000014 |lr_p 0.000018 | batch 10304 | examples/s: 113.6 | loss: 0.69125 | time e
lapsed: 02h58m08s | time left: 02h13m58s
Training
epoch 25 | lr 0.000011 |lr_p 0.000016 | batch 900 | examples/s: 105.8 | loss: 0.38796 | time e
lapsed: 02h59m23s | time left: 02h12m43s
epoch 25 | lr 0.000011 |lr_p 0.000016 | batch 2900 | examples/s: 110.9 | loss: 0.71663 | time e
lapsed: 03h00m36s | time left: 02h11m26s
epoch 25 | lr 0.000011 |lr_p 0.000016 | batch 4900 | examples/s: 113.7 | loss: 0.99163 | time e
lapsed: 03h01m47s | time left: 02h10m08s
epoch 25 | lr 0.000011 |lr_p 0.000016 | batch 6900 | examples/s: 114.5 | loss: 0.27221 | time e
lapsed: 03h02m59s | time left: 02h08m50s
epoch 25 | lr 0.000011 |lr_p 0.000016 | batch 8900 | examples/s: 113.2 | loss: 0.46326 | time e
lapsed: 03h04m10s | time left: 02h07m32s
epoch 25 | lr 0.000011 |lr_p 0.000016 | batch 10900 | examples/s: 112.6 | loss: 0.61697 | time e
lapsed: 03h05m22s | time left: 02h06m15s
Training
epoch 26 | lr 0.000009 |lr_p 0.000014 | batch 1496 | examples/s: 113.6 | loss: 0.81528 | time e
lapsed: 03h06m37s | time left: 02h05m00s
epoch 26 | lr 0.000009 |lr_p 0.000014 | batch 3496 | examples/s: 109.1 | loss: 0.33061 | time e
lapsed: 03h07m50s | time left: 02h03m43s
epoch 26 | lr 0.000009 |lr_p 0.000014 | batch 5496 | examples/s: 109.8 | loss: 0.42606 | time e
lapsed: 03h09m02s | time left: 02h02m26s

```

```

epoch 26 | lr 0.000009 |lr_p 0.000014 | batch 7496 | examples/s: 114.3 | loss: 0.69500 | time e
lapsed: 03h10m14s | time left: 02h01m09s
epoch 26 | lr 0.000009 |lr_p 0.000014 | batch 9496 | examples/s: 105.0 | loss: 0.44754 | time e
lapsed: 03h11m26s | time left: 01h59m52s
Training
epoch 27 | lr 0.000007 |lr_p 0.000012 | batch 92 | examples/s: 112.9 | loss: 0.59403 | time e
lapsed: 03h12m40s | time left: 01h58m36s
epoch 27 | lr 0.000007 |lr_p 0.000012 | batch 2092 | examples/s: 111.4 | loss: 0.41021 | time e
lapsed: 03h13m54s | time left: 01h57m20s
epoch 27 | lr 0.000007 |lr_p 0.000012 | batch 4092 | examples/s: 111.3 | loss: 0.71705 | time e
lapsed: 03h15m06s | time left: 01h56m04s
epoch 27 | lr 0.000007 |lr_p 0.000012 | batch 6092 | examples/s: 109.9 | loss: 0.60577 | time e
lapsed: 03h16m19s | time left: 01h54m47s
epoch 27 | lr 0.000007 |lr_p 0.000012 | batch 8092 | examples/s: 113.0 | loss: 0.41555 | time e
lapsed: 03h17m31s | time left: 01h53m30s
epoch 27 | lr 0.000007 |lr_p 0.000012 | batch 10092 | examples/s: 114.9 | loss: 0.14415 | time e
lapsed: 03h18m43s | time left: 01h52m14s
Training
epoch 28 | lr 0.000006 |lr_p 0.000011 | batch 688 | examples/s: 113.8 | loss: 0.65281 | time e
lapsed: 03h19m58s | time left: 01h50m59s
epoch 28 | lr 0.000006 |lr_p 0.000011 | batch 2688 | examples/s: 89.0 | loss: 0.57737 | time e
lapsed: 03h21m11s | time left: 01h49m42s
epoch 28 | lr 0.000006 |lr_p 0.000011 | batch 4688 | examples/s: 112.8 | loss: 0.60705 | time e
lapsed: 03h22m22s | time left: 01h48m26s
epoch 28 | lr 0.000006 |lr_p 0.000011 | batch 6688 | examples/s: 108.0 | loss: 0.34504 | time e
lapsed: 03h23m34s | time left: 01h47m09s
epoch 28 | lr 0.000006 |lr_p 0.000011 | batch 8688 | examples/s: 116.2 | loss: 0.59514 | time e
lapsed: 03h24m45s | time left: 01h45m52s
epoch 28 | lr 0.000006 |lr_p 0.000011 | batch 10688 | examples/s: 111.1 | loss: 0.27837 | time e
lapsed: 03h25m58s | time left: 01h44m36s
Training
epoch 29 | lr 0.000005 |lr_p 0.000010 | batch 1284 | examples/s: 106.3 | loss: 0.64420 | time e
lapsed: 03h27m12s | time left: 01h43m21s
epoch 29 | lr 0.000005 |lr_p 0.000010 | batch 3284 | examples/s: 91.3 | loss: 0.20391 | time e
lapsed: 03h28m26s | time left: 01h42m05s
epoch 29 | lr 0.000005 |lr_p 0.000010 | batch 5284 | examples/s: 109.8 | loss: 0.65798 | time e
lapsed: 03h29m39s | time left: 01h40m49s
epoch 29 | lr 0.000005 |lr_p 0.000010 | batch 7284 | examples/s: 115.7 | loss: 0.32823 | time e
lapsed: 03h30m50s | time left: 01h39m33s
epoch 29 | lr 0.000005 |lr_p 0.000010 | batch 9284 | examples/s: 108.7 | loss: 0.55189 | time e
lapsed: 03h32m02s | time left: 01h38m17s
epoch 29 | lr 0.000005 |lr_p 0.000010 | batch 11284 | examples/s: 111.6 | loss: 0.42040 | time e
lapsed: 03h33m14s | time left: 01h37m01s
Training
epoch 30 | lr 0.000090 |lr_p 0.000090 | batch 1880 | examples/s: 103.5 | loss: 0.59631 | time e
lapsed: 03h34m32s | time left: 01h35m47s
epoch 30 | lr 0.000090 |lr_p 0.000090 | batch 3880 | examples/s: 102.2 | loss: 0.45386 | time e
lapsed: 03h35m50s | time left: 01h34m34s
epoch 30 | lr 0.000090 |lr_p 0.000090 | batch 5880 | examples/s: 102.2 | loss: 0.52599 | time e
lapsed: 03h37m09s | time left: 01h33m21s
epoch 30 | lr 0.000090 |lr_p 0.000090 | batch 7880 | examples/s: 105.8 | loss: 0.90160 | time e
lapsed: 03h38m27s | time left: 01h32m07s
epoch 30 | lr 0.000090 |lr_p 0.000090 | batch 9880 | examples/s: 99.8 | loss: 0.30211 | time e
lapsed: 03h39m45s | time left: 01h30m53s
Training
epoch 31 | lr 0.000090 |lr_p 0.000090 | batch 476 | examples/s: 103.2 | loss: 0.63951 | time e
lapsed: 03h41m06s | time left: 01h29m41s
epoch 31 | lr 0.000090 |lr_p 0.000090 | batch 2476 | examples/s: 102.5 | loss: 0.23640 | time e
lapsed: 03h42m24s | time left: 01h28m27s
epoch 31 | lr 0.000090 |lr_p 0.000090 | batch 4476 | examples/s: 107.8 | loss: 0.54902 | time e
lapsed: 03h43m41s | time left: 01h27m13s
epoch 31 | lr 0.000090 |lr_p 0.000090 | batch 6476 | examples/s: 102.4 | loss: 0.39028 | time e
lapsed: 03h44m59s | time left: 01h25m59s
epoch 31 | lr 0.000090 |lr_p 0.000090 | batch 8476 | examples/s: 99.4 | loss: 0.52520 | time e
lapsed: 03h46m16s | time left: 01h24m45s
epoch 31 | lr 0.000090 |lr_p 0.000090 | batch 10476 | examples/s: 105.8 | loss: 0.45191 | time e
lapsed: 03h47m35s | time left: 01h23m31s
Training
epoch 32 | lr 0.000089 |lr_p 0.000089 | batch 1072 | examples/s: 104.6 | loss: 0.48635 | time e
lapsed: 03h48m57s | time left: 01h22m19s
epoch 32 | lr 0.000089 |lr_p 0.000089 | batch 3072 | examples/s: 102.0 | loss: 0.29880 | time e
lapsed: 03h50m17s | time left: 01h21m06s
epoch 32 | lr 0.000089 |lr_p 0.000089 | batch 5072 | examples/s: 99.3 | loss: 0.34316 | time e
lapsed: 03h51m37s | time left: 01h19m52s
epoch 32 | lr 0.000089 |lr_p 0.000089 | batch 7072 | examples/s: 103.4 | loss: 0.36289 | time e
lapsed: 03h52m56s | time left: 01h18m38s
epoch 32 | lr 0.000089 |lr_p 0.000089 | batch 9072 | examples/s: 99.6 | loss: 0.31858 | time e
lapsed: 03h54m15s | time left: 01h17m25s
epoch 32 | lr 0.000089 |lr_p 0.000089 | batch 11072 | examples/s: 98.9 | loss: 0.69952 | time e
lapsed: 03h55m35s | time left: 01h16m11s

```

Training

epoch 33 lr 0.000088 lr_p 0.000088 batch	1668 examples/s: 92.0 loss: 0.69925 time e
lapsed: 03h56m57s time left: 01h14m58s	
epoch 33 lr 0.000088 lr_p 0.000088 batch	3668 examples/s: 101.9 loss: 0.66058 time e
lapsed: 03h58m16s time left: 01h13m44s	
epoch 33 lr 0.000088 lr_p 0.000088 batch	5668 examples/s: 102.5 loss: 0.37866 time e
lapsed: 03h59m35s time left: 01h12m30s	
epoch 33 lr 0.000088 lr_p 0.000088 batch	7668 examples/s: 101.2 loss: 0.55513 time e
lapsed: 04h00m54s time left: 01h11m16s	
epoch 33 lr 0.000088 lr_p 0.000088 batch	9668 examples/s: 95.8 loss: 0.34259 time e
lapsed: 04h02m13s time left: 01h10m02s	

Training

epoch 34 lr 0.000087 lr_p 0.000087 batch	264 examples/s: 81.7 loss: 0.60197 time e
lapsed: 04h03m35s time left: 01h08m48s	
epoch 34 lr 0.000087 lr_p 0.000087 batch	2264 examples/s: 104.8 loss: 0.24841 time e
lapsed: 04h04m54s time left: 01h07m34s	
epoch 34 lr 0.000087 lr_p 0.000087 batch	4264 examples/s: 104.3 loss: 0.34349 time e
lapsed: 04h06m14s time left: 01h06m19s	
epoch 34 lr 0.000087 lr_p 0.000087 batch	6264 examples/s: 100.3 loss: 0.95180 time e
lapsed: 04h07m33s time left: 01h05m05s	
epoch 34 lr 0.000087 lr_p 0.000087 batch	8264 examples/s: 103.1 loss: 0.18813 time e
lapsed: 04h08m52s time left: 01h03m51s	
epoch 34 lr 0.000087 lr_p 0.000087 batch	10264 examples/s: 101.1 loss: 0.22310 time e
lapsed: 04h10m11s time left: 01h02m36s	

Training

epoch 35 lr 0.000085 lr_p 0.000085 batch	860 examples/s: 101.5 loss: 0.68612 time e
lapsed: 04h11m33s time left: 01h01m22s	
epoch 35 lr 0.000085 lr_p 0.000085 batch	2860 examples/s: 98.4 loss: 0.34472 time e
lapsed: 04h12m53s time left: 01h00m08s	
epoch 35 lr 0.000085 lr_p 0.000085 batch	4860 examples/s: 99.5 loss: 0.38941 time e
lapsed: 04h14m13s time left: 00h58m54s	
epoch 35 lr 0.000085 lr_p 0.000085 batch	6860 examples/s: 93.5 loss: 0.56871 time e
lapsed: 04h15m33s time left: 00h57m39s	
epoch 35 lr 0.000085 lr_p 0.000085 batch	8860 examples/s: 104.9 loss: 0.28711 time e
lapsed: 04h16m52s time left: 00h56m24s	
epoch 35 lr 0.000085 lr_p 0.000085 batch	10860 examples/s: 81.8 loss: 0.62001 time e
lapsed: 04h18m11s time left: 00h55m09s	

Training

epoch 36 lr 0.000082 lr_p 0.000083 batch	1456 examples/s: 96.0 loss: 0.34747 time e
lapsed: 04h19m33s time left: 00h53m55s	
epoch 36 lr 0.000082 lr_p 0.000083 batch	3456 examples/s: 98.6 loss: 0.61646 time e
lapsed: 04h20m53s time left: 00h52m41s	
epoch 36 lr 0.000082 lr_p 0.000083 batch	5456 examples/s: 99.5 loss: 1.02825 time e
lapsed: 04h22m13s time left: 00h51m26s	
epoch 36 lr 0.000082 lr_p 0.000083 batch	7456 examples/s: 101.2 loss: 0.28100 time e
lapsed: 04h23m33s time left: 00h50m11s	
epoch 36 lr 0.000082 lr_p 0.000083 batch	9456 examples/s: 99.0 loss: 0.41859 time e
lapsed: 04h24m53s time left: 00h48m56s	

Training

epoch 37 lr 0.000080 lr_p 0.000080 batch	52 examples/s: 100.2 loss: 0.64991 time e
lapsed: 04h26m15s time left: 00h47m41s	
epoch 37 lr 0.000080 lr_p 0.000080 batch	2052 examples/s: 104.9 loss: 0.56187 time e
lapsed: 04h27m34s time left: 00h46m26s	
epoch 37 lr 0.000080 lr_p 0.000080 batch	4052 examples/s: 98.1 loss: 0.21505 time e
lapsed: 04h28m54s time left: 00h45m11s	
epoch 37 lr 0.000080 lr_p 0.000080 batch	6052 examples/s: 100.9 loss: 0.24201 time e
lapsed: 04h30m14s time left: 00h43m56s	
epoch 37 lr 0.000080 lr_p 0.000080 batch	8052 examples/s: 98.8 loss: 0.47542 time e
lapsed: 04h31m32s time left: 00h42m41s	
epoch 37 lr 0.000080 lr_p 0.000080 batch	10052 examples/s: 105.6 loss: 0.56635 time e
lapsed: 04h32m51s time left: 00h41m26s	

Training

epoch 38 lr 0.000077 lr_p 0.000078 batch	648 examples/s: 102.5 loss: 0.51054 time e
lapsed: 04h34m12s time left: 00h40m11s	
epoch 38 lr 0.000077 lr_p 0.000078 batch	2648 examples/s: 94.6 loss: 0.53858 time e
lapsed: 04h35m32s time left: 00h38m55s	
epoch 38 lr 0.000077 lr_p 0.000078 batch	4648 examples/s: 98.9 loss: 0.59242 time e
lapsed: 04h36m53s time left: 00h37m40s	
epoch 38 lr 0.000077 lr_p 0.000078 batch	6648 examples/s: 100.6 loss: 0.31681 time e
lapsed: 04h38m13s time left: 00h36m25s	
epoch 38 lr 0.000077 lr_p 0.000078 batch	8648 examples/s: 103.5 loss: 0.36978 time e
lapsed: 04h39m32s time left: 00h35m09s	
epoch 38 lr 0.000077 lr_p 0.000078 batch	10648 examples/s: 99.7 loss: 0.47957 time e
lapsed: 04h40m52s time left: 00h33m54s	

Training

epoch 39 lr 0.000074 lr_p 0.000074 batch	1244 examples/s: 100.5 loss: 0.48269 time e
lapsed: 04h42m14s time left: 00h32m39s	
epoch 39 lr 0.000074 lr_p 0.000074 batch	3244 examples/s: 104.4 loss: 0.33208 time e
lapsed: 04h43m33s time left: 00h31m23s	
epoch 39 lr 0.000074 lr_p 0.000074 batch	5244 examples/s: 104.9 loss: 0.27144 time e

lapsed: 04h44m52s | time left: 00h30m08s
epoch 39 | lr 0.000074 | lr_p 0.000074 | batch 7244 | examples/s: 105.3 | loss: 0.87697 | time e
lapsed: 04h46m11s | time left: 00h28m52s
epoch 39 | lr 0.000074 | lr_p 0.000074 | batch 9244 | examples/s: 105.4 | loss: 0.44309 | time e
lapsed: 04h47m29s | time left: 00h27m36s
epoch 39 | lr 0.000074 | lr_p 0.000074 | batch 11244 | examples/s: 105.1 | loss: 0.28501 | time e
lapsed: 04h48m50s | time left: 00h26m21s
Training
epoch 40 | lr 0.000070 | lr_p 0.000071 | batch 1840 | examples/s: 104.9 | loss: 1.26091 | time e
lapsed: 04h50m12s | time left: 00h25m05s
epoch 40 | lr 0.000070 | lr_p 0.000071 | batch 3840 | examples/s: 94.4 | loss: 0.35300 | time e
lapsed: 04h51m31s | time left: 00h23m49s
epoch 40 | lr 0.000070 | lr_p 0.000071 | batch 5840 | examples/s: 96.8 | loss: 0.38632 | time e
lapsed: 04h52m50s | time left: 00h22m33s
epoch 40 | lr 0.000070 | lr_p 0.000071 | batch 7840 | examples/s: 103.5 | loss: 0.21808 | time e
lapsed: 04h54m09s | time left: 00h21m18s
epoch 40 | lr 0.000070 | lr_p 0.000071 | batch 9840 | examples/s: 84.9 | loss: 0.30779 | time e
lapsed: 04h55m29s | time left: 00h20m02s
Training
epoch 41 | lr 0.000066 | lr_p 0.000068 | batch 436 | examples/s: 93.3 | loss: 0.48145 | time e
lapsed: 04h56m50s | time left: 00h18m46s
epoch 41 | lr 0.000066 | lr_p 0.000068 | batch 2436 | examples/s: 100.3 | loss: 0.67462 | time e
lapsed: 04h58m09s | time left: 00h17m30s
epoch 41 | lr 0.000066 | lr_p 0.000068 | batch 4436 | examples/s: 103.6 | loss: 0.62339 | time e
lapsed: 04h59m28s | time left: 00h16m14s
epoch 41 | lr 0.000066 | lr_p 0.000068 | batch 6436 | examples/s: 102.3 | loss: 0.49116 | time e
lapsed: 05h00m47s | time left: 00h14m58s
epoch 41 | lr 0.000066 | lr_p 0.000068 | batch 8436 | examples/s: 115.6 | loss: 0.33865 | time e
lapsed: 05h02m00s | time left: 00h13m42s
epoch 41 | lr 0.000066 | lr_p 0.000068 | batch 10436 | examples/s: 114.7 | loss: 0.25293 | time e
lapsed: 05h03m10s | time left: 00h12m25s
Training
epoch 42 | lr 0.000062 | lr_p 0.000064 | batch 1032 | examples/s: 116.0 | loss: 0.40150 | time e
lapsed: 05h04m24s | time left: 00h11m09s
epoch 42 | lr 0.000062 | lr_p 0.000064 | batch 3032 | examples/s: 112.3 | loss: 0.54956 | time e
lapsed: 05h05m37s | time left: 00h09m53s
epoch 42 | lr 0.000062 | lr_p 0.000064 | batch 5032 | examples/s: 105.6 | loss: 0.55137 | time e
lapsed: 05h06m50s | time left: 00h08m37s
epoch 42 | lr 0.000062 | lr_p 0.000064 | batch 7032 | examples/s: 114.4 | loss: 0.29713 | time e
lapsed: 05h08m02s | time left: 00h07m21s
epoch 42 | lr 0.000062 | lr_p 0.000064 | batch 9032 | examples/s: 114.0 | loss: 0.45576 | time e
lapsed: 05h09m14s | time left: 00h06m04s
epoch 42 | lr 0.000062 | lr_p 0.000064 | batch 11032 | examples/s: 112.6 | loss: 0.66483 | time e
lapsed: 05h10m26s | time left: 00h04m48s
Training
epoch 43 | lr 0.000058 | lr_p 0.000060 | batch 1628 | examples/s: 113.4 | loss: 0.50200 | time e
lapsed: 05h11m41s | time left: 00h03m32s
epoch 43 | lr 0.000058 | lr_p 0.000060 | batch 3628 | examples/s: 102.4 | loss: 0.68082 | time e
lapsed: 05h12m53s | time left: 00h02m16s
epoch 43 | lr 0.000058 | lr_p 0.000060 | batch 5628 | examples/s: 108.0 | loss: 0.57775 | time e
lapsed: 05h14m05s | time left: 00h01m00s
epoch 43 | lr 0.000058 | lr_p 0.000060 | batch 7628 | examples/s: 113.8 | loss: 0.31813 | time e
lapsed: 05h15m17s | time left: -1h59m45s
epoch 43 | lr 0.000058 | lr_p 0.000060 | batch 9628 | examples/s: 115.1 | loss: 0.35753 | time e
lapsed: 05h16m28s | time left: -1h58m29s
Training
epoch 44 | lr 0.000054 | lr_p 0.000056 | batch 224 | examples/s: 106.0 | loss: 0.64282 | time e
lapsed: 05h17m43s | time left: -1h57m13s
epoch 44 | lr 0.000054 | lr_p 0.000056 | batch 2224 | examples/s: 113.4 | loss: 0.40013 | time e
lapsed: 05h18m56s | time left: -1h55m57s
epoch 44 | lr 0.000054 | lr_p 0.000056 | batch 4224 | examples/s: 113.2 | loss: 0.71669 | time e
lapsed: 05h20m09s | time left: -1h54m42s
epoch 44 | lr 0.000054 | lr_p 0.000056 | batch 6224 | examples/s: 110.8 | loss: 0.43887 | time e
lapsed: 05h21m21s | time left: -1h53m26s
epoch 44 | lr 0.000054 | lr_p 0.000056 | batch 8224 | examples/s: 114.6 | loss: 0.67291 | time e
lapsed: 05h22m33s | time left: -1h52m10s
epoch 44 | lr 0.000054 | lr_p 0.000056 | batch 10224 | examples/s: 113.1 | loss: 0.46929 | time e
lapsed: 05h23m45s | time left: -1h50m54s
Training
epoch 45 | lr 0.000050 | lr_p 0.000052 | batch 820 | examples/s: 101.8 | loss: 0.32622 | time e
lapsed: 05h25m00s | time left: -1h49m38s
epoch 45 | lr 0.000050 | lr_p 0.000052 | batch 2820 | examples/s: 105.3 | loss: 0.63897 | time e
lapsed: 05h26m13s | time left: -1h48m23s
epoch 45 | lr 0.000050 | lr_p 0.000052 | batch 4820 | examples/s: 113.9 | loss: 0.39574 | time e
lapsed: 05h27m25s | time left: -1h47m07s
epoch 45 | lr 0.000050 | lr_p 0.000052 | batch 6820 | examples/s: 117.0 | loss: 0.33405 | time e
lapsed: 05h28m37s | time left: -1h45m51s
epoch 45 | lr 0.000050 | lr_p 0.000052 | batch 8820 | examples/s: 110.3 | loss: 0.86956 | time e
lapsed: 05h29m49s | time left: -1h44m35s
epoch 45 | lr 0.000050 | lr_p 0.000052 | batch 10820 | examples/s: 110.9 | loss: 0.64461 | time e

```

lapsed: 05h31m01s | time left: -1h43m20s
Training
epoch 46 | lr 0.000045 |lr_p 0.000048 | batch 1416 | examples/s: 116.4 | loss: 0.38747 | time e
lapsed: 05h32m16s | time left: -1h42m04s
epoch 46 | lr 0.000045 |lr_p 0.000048 | batch 3416 | examples/s: 106.4 | loss: 0.24114 | time e
lapsed: 05h33m28s | time left: -1h40m48s
epoch 46 | lr 0.000045 |lr_p 0.000048 | batch 5416 | examples/s: 113.2 | loss: 0.68167 | time e
lapsed: 05h34m41s | time left: -1h39m33s
epoch 46 | lr 0.000045 |lr_p 0.000048 | batch 7416 | examples/s: 117.3 | loss: 0.28477 | time e
lapsed: 05h35m53s | time left: -1h38m17s
epoch 46 | lr 0.000045 |lr_p 0.000048 | batch 9416 | examples/s: 113.6 | loss: 0.43869 | time e
lapsed: 05h37m05s | time left: -1h37m02s
Training
epoch 47 | lr 0.000041 |lr_p 0.000044 | batch 12 | examples/s: 103.2 | loss: 0.27357 | time e
lapsed: 05h38m20s | time left: -1h35m46s
epoch 47 | lr 0.000041 |lr_p 0.000044 | batch 2012 | examples/s: 109.0 | loss: 0.32599 | time e
lapsed: 05h39m33s | time left: -1h34m31s
epoch 47 | lr 0.000041 |lr_p 0.000044 | batch 4012 | examples/s: 115.9 | loss: 0.54442 | time e
lapsed: 05h40m45s | time left: -1h33m15s
epoch 47 | lr 0.000041 |lr_p 0.000044 | batch 6012 | examples/s: 116.0 | loss: 0.41749 | time e
lapsed: 05h41m57s | time left: -1h32m00s
epoch 47 | lr 0.000041 |lr_p 0.000044 | batch 8012 | examples/s: 115.4 | loss: 0.56326 | time e
lapsed: 05h43m09s | time left: -1h30m44s
epoch 47 | lr 0.000041 |lr_p 0.000044 | batch 10012 | examples/s: 109.6 | loss: 0.22976 | time e
lapsed: 05h44m21s | time left: -1h29m29s
Training
epoch 48 | lr 0.000037 |lr_p 0.000040 | batch 608 | examples/s: 116.7 | loss: 0.30489 | time e
lapsed: 05h45m35s | time left: -1h28m13s
epoch 48 | lr 0.000037 |lr_p 0.000040 | batch 2608 | examples/s: 112.8 | loss: 0.28081 | time e
lapsed: 05h46m48s | time left: -1h26m58s
epoch 48 | lr 0.000037 |lr_p 0.000040 | batch 4608 | examples/s: 112.4 | loss: 0.30178 | time e
lapsed: 05h47m59s | time left: -1h25m43s
epoch 48 | lr 0.000037 |lr_p 0.000040 | batch 6608 | examples/s: 100.4 | loss: 0.49426 | time e
lapsed: 05h49m12s | time left: -1h24m27s
epoch 48 | lr 0.000037 |lr_p 0.000040 | batch 8608 | examples/s: 114.6 | loss: 0.27966 | time e
lapsed: 05h50m23s | time left: -1h23m12s
epoch 48 | lr 0.000037 |lr_p 0.000040 | batch 10608 | examples/s: 111.1 | loss: 0.22168 | time e
lapsed: 05h51m35s | time left: -1h21m57s
Training
epoch 49 | lr 0.000033 |lr_p 0.000036 | batch 1204 | examples/s: 104.9 | loss: 0.47220 | time e
lapsed: 05h52m49s | time left: -1h20m42s
epoch 49 | lr 0.000033 |lr_p 0.000036 | batch 3204 | examples/s: 109.3 | loss: 0.79936 | time e
lapsed: 05h54m02s | time left: -1h19m26s
epoch 49 | lr 0.000033 |lr_p 0.000036 | batch 5204 | examples/s: 111.9 | loss: 0.21000 | time e
lapsed: 05h55m14s | time left: -1h18m11s
epoch 49 | lr 0.000033 |lr_p 0.000036 | batch 7204 | examples/s: 106.9 | loss: 0.38305 | time e
lapsed: 05h56m26s | time left: -1h16m56s
epoch 49 | lr 0.000033 |lr_p 0.000036 | batch 9204 | examples/s: 107.0 | loss: 0.31274 | time e
lapsed: 05h57m38s | time left: -1h15m41s
epoch 49 | lr 0.000033 |lr_p 0.000036 | batch 11204 | examples/s: 113.4 | loss: 0.47945 | time e
lapsed: 05h58m50s | time left: -1h14m26s

```

```
In [ ]: # pip install 'git+https://github.com/saadnaeem-dev/pytorch-linear-warmup-cosine-annealing-warm-re
```

```
In [ ]: import os
os.getcwd()
```

```
Out[ ]: '/mnt/workspace/sunqiao/mymono'
```

```
In [ ]: pwd
```

```
Out[ ]: '/mnt/workspace/sunqiao/mymono'
```