

# Imports

```
In [ ]: import numpy as np
import pandas as pd
import torch
import torch.nn as nn
from torch.cuda.amp import GradScaler, autocast
import matplotlib.pyplot as plt
from tqdm.auto import tqdm
from pathlib import Path
import cv2 as cv
from PIL import Image
import segmentation_models_pytorch as smp
from sklearn.model_selection import train_test_split
import albumentations as A
from albumentations.pytorch import ToTensorV2
from torchmetrics.image import StructuralSimilarityIndexMeasure as SSIM
from torchmetrics.regression import MeanSquaredError as MSE
from torchmetrics.collections import MetricCollection
import gc
from torchvision.transforms import Normalize
```

```
In [ ]: torch.cuda.device_count()
```

```
Out[ ]: 2
```

```
In [ ]: for i in range(torch.cuda.device_count()):
    print(torch.cuda.get_device_name(i))
```

```
NVIDIA GeForce RTX 2080 Ti
NVIDIA GeForce RTX 2080 Ti
```

```
In [ ]: # device =
# torch.device('cuda:0'), torch.device('cuda:1'), torch.device('cuda:3')
```

# Data & Visualization

**Note: brighter the area, farther it is in the image, brighter => more depth**

```
In [ ]: train_csv = Path('./nyu-depth-v2/nyu_data/data/nyu2_train.csv')
train_ims_path = Path('./nyu-depth-v2/nyu_data/data/nyu2_train')
base_path = Path('./nyu-depth-v2/nyu_data')
```

```
In [ ]: df = pd.read_csv(train_csv, header=None)
df[0] = df[0].map(lambda x:base_path/x)
df[1] = df[1].map(lambda x:base_path/x)
df.head()
```

Out[ ]:

0

1

```
0 nyu-depth-v2\nyu_data\data\nyu2_train\living_r... nyu-depth-v2\nyu_data\data\nyu2_train\living_r...
1 nyu-depth-v2\nyu_data\data\nyu2_train\living_r... nyu-depth-v2\nyu_data\data\nyu2_train\living_r...
2 nyu-depth-v2\nyu_data\data\nyu2_train\living_r... nyu-depth-v2\nyu_data\data\nyu2_train\living_r...
3 nyu-depth-v2\nyu_data\data\nyu2_train\living_r... nyu-depth-v2\nyu_data\data\nyu2_train\living_r...
4 nyu-depth-v2\nyu_data\data\nyu2_train\living_r... nyu-depth-v2\nyu_data\data\nyu2_train\living_r...
```

In [ ]:

```
train_df, val_df = train_test_split(df, test_size=0.1, shuffle=True)
val_df, test_df = train_test_split(val_df, test_size=0.1, shuffle=True)
train_df.reset_index(drop=True, inplace=True)
val_df.reset_index(drop=True, inplace=True)
test_df.reset_index(drop=True, inplace=True)
len(train_df), len(val_df), len(test_df)
```

Out[ ]:

(45619, 4562, 507)

In [ ]:

```
def colored_depthmap(depth, d_min=None, d_max=None, cmap=plt.cm.inferno):
    if d_min is None:
        d_min = np.min(depth)
    if d_max is None:
        d_max = np.max(depth)
    depth_relative = (depth - d_min) / (d_max - d_min)
    return 255 * cmap(depth_relative)[:, :, :3] # H, W, C

def merge_into_row(input, depth_target):
    input = np.array(input)
    depth_target = np.squeeze(np.array(depth_target))

    d_min = np.min(depth_target)
    d_max = np.max(depth_target)
    depth_target_col = colored_depthmap(depth_target, d_min, d_max)
    img_merge = np.hstack([input, depth_target_col])

    return img_merge

plt.figure(figsize=(15, 6))
for i, idx in enumerate(np.random.randint(0, len(df), (16,))):
    ax = plt.subplot(4, 4, i + 1)
    image = Image.open(df.iloc[idx, 0]).convert('RGB')
    mask = Image.open(df.iloc[idx, 1]).convert('L')
    image_viz = merge_into_row(
        image, mask
    )
    plt.imshow(image_viz.astype("uint8"))
    plt.axis("off")
```



## Dataset

```
In [ ]: sample_tfms = [
    A.HorizontalFlip(),
    A.GaussNoise(p=0.2),
    A.OneOf([
        A.MotionBlur(p=.3),
        A.MedianBlur(blur_limit=3, p=0.3),
        A.Blur(blur_limit=3, p=0.5),
    ], p=0.3),
    A.RGBShift(),
    A.RandomBrightnessContrast(),
    A.RandomResizedCrop(384, 384),
    A.ColorJitter(),
    A.ShiftScaleRotate(shift_limit=0.1, scale_limit=0.3, rotate_limit=45, p=0.5),
    A.HueSaturationValue(p=0.3),
]
train_tfms = A.Compose([
    *sample_tfms,
    A.Resize(224, 224),
    A.Normalize(always_apply=True),
    ToTensorV2()
])
valid_tfms = A.Compose([
    A.Resize(224, 224),
    A.Normalize(always_apply=True),
    ToTensorV2()
])
```

```
In [ ]: class Dataset:
    def __init__(self, df, tfms):
        self.df = df
        self.tfms = tfms
    def open_im(self, p, gray=False):
        im = cv.imread(str(p))
        im = cv.cvtColor(im, cv.COLOR_BGR2GRAY if gray else cv.COLOR_BGR2RGB)
        return im

    def __len__(self):
        return len(self.df)

    def __getitem__(self, idx):
        s = self.df.iloc[idx, :]
        im, dp = s[0], s[1]
        im, dp = self.open_im(im), self.open_im(dp, True)
```

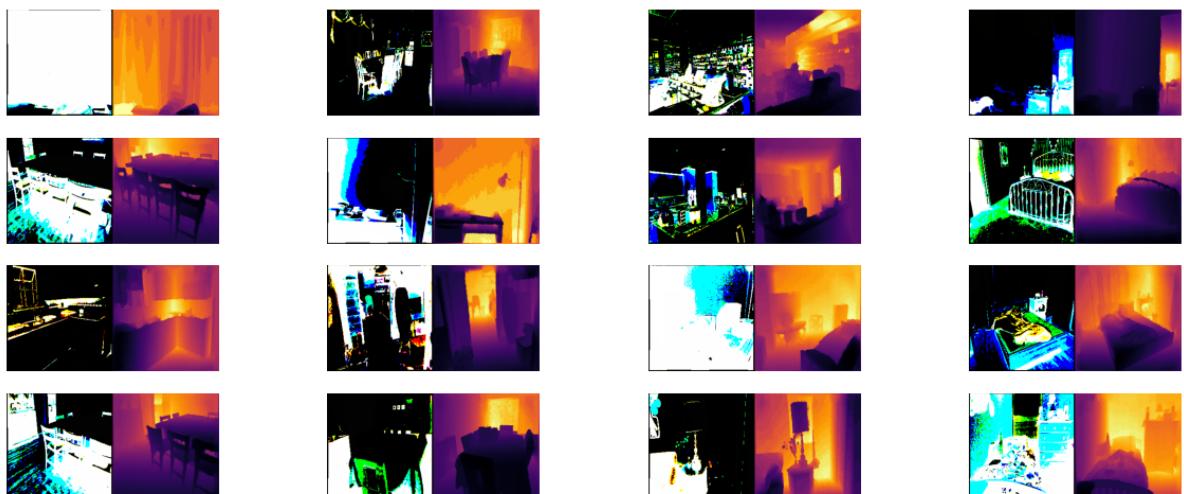
```
augs = self.tfms(image=im, mask=dp)
im, dp = augs['image'], augs['mask'] / 255.
return im, dp.unsqueeze(0)
```

In [ ]: train\_ds = Dataset(train\_df, train\_tfms)  
val\_ds = Dataset(val\_df, valid\_tfms)  
test\_ds = Dataset(test\_df, valid\_tfms)  
len(train\_ds), len(val\_ds), len(test\_ds)

Out[ ]: (45619, 4562, 507)

## Transforms Visualization

In [ ]: randoms = [test\_ds[idx] for idx in range(16)]
plt.figure(figsize=(15, 6))
for i, (img, mask) in enumerate(randoms):
 ax = plt.subplot(4, 4, i + 1)
 img, mask = img.permute(1, 2, 0).numpy(), mask.permute(1, 2, 0).numpy()
 mask = colored\_depthmap(np.squeeze(mask))
 image\_viz = np.hstack([img, mask])
 plt.imshow(image\_viz.astype("uint8"))
 plt.axis("off")



In [ ]: train\_dl = torch.utils.data.DataLoader(train\_ds, shuffle=True, batch\_size=64)
val\_dl = torch.utils.data.DataLoader(val\_ds, shuffle=False, batch\_size=64)
test\_dl = torch.utils.data.DataLoader(test\_ds, shuffle=True, batch\_size=4)
len(train\_dl), len(val\_dl), len(test\_dl)

Out[ ]: (713, 72, 127)

## Model

In [ ]: class UNet(nn.Module):
 def \_\_init\_\_(self):
 super().\_\_init\_\_()
 self.model = smp.UnetPlusPlus(
 encoder\_name='resnext50\_32x4d',
 in\_channels=3,
 classes=1
 )

 def trainable\_encoder(self, trainable=True):

```

        for p in self.model.encoder.parameters():
            p.requires_grad = trainable

    def forward(self, x):
        return self.model(x)

    def _num_params(self,):
        return sum([p.numel() for p in self.model.parameters() if p.requires_grad])

```

# Training

## Metrics:

### Structural Similarity Index (SSIM)

The structural similarity index (SSIM) is a perceptual metric that takes into account the structural similarity of the two images. SSIM is calculated by comparing the local patterns of the two images, taking into account the luminance, contrast, and structure of the images.

The formula for SSIM is:

$$SSIM = \frac{(2\mu_x\mu_y + c_1)(2\sigma_{xy} + c_2)}{(\mu_x^2 + \mu_y^2 + c_1)(\sigma_x^2 + \sigma_y^2 + c_2)}$$

where  $\mu_x$  and  $\mu_y$  are the means of the two images,  $\sigma_x^2$  and  $\sigma_y^2$  are the variances of the two images, and  $\sigma_{xy}$  is the covariance of the two images.  $c_1$  and  $c_2$  are constants that are used to stabilize the calculation of SSIM.

SSIM is more robust to noise and small changes than MSE, but it is also more computationally expensive to calculate.

In general, SSIM is preferred over MSE for image quality assessment because it provides a more accurate measure of how humans perceive the similarity of two images. However, MSE is still a useful metric, especially when speed is important.

```

In [ ]: class UnNormalize(Normalize):
    def __init__(self, *args, **kwargs):
        mean=(0.485, 0.456, 0.406)
        std=(0.229, 0.224, 0.225)
        new_mean = [-m/s for m, s in zip(mean, std)]
        new_std = [1/s for s in std]
        super().__init__(new_mean, new_std, *args, **kwargs)

    @torch.no_grad()
    def plot_vals(imgs, preds, targets, n=4, figsize=(6, 2), title=''):
        plt.figure(figsize=figsize, dpi=150)
        r = 2 if n == 4 else 8
        c = 2
        for i, idx in enumerate(np.random.randint(0, imgs.size(0), (n,))):
            ax = plt.subplot(r, c, i + 1)
            img, pred, gt = imgs[idx], preds[idx], targets[idx]
            img = UnNormalize()(img)*255.
            img, pred, gt = img.permute(1, 2, 0).numpy(), pred.permute(1, 2, 0).numpy(), gt.permute(1, 2, 0).numpy()
            pred = colored_depthmap(np.squeeze(pred))
            gt = colored_depthmap(np.squeeze(gt))
            image_viz = np.hstack([img, pred, gt])

```

```

plt.imshow(image_viz.astype("uint8"))
plt.axis("off")
title = f'{title}\nimage/target/prediction' if len(title)!=0 else 'image/target'
plt.suptitle(title)
plt.show()

```

```

In [ ]: epochs = 5
freeze_epochs = 2
lr = 1e-3

device = torch.device('cuda:1')

metrics = MetricCollection([
    SSIM(data_range=(0, 1)),
    MSE()
]).to(device)
train_metrics = metrics.clone()
val_metrics = metrics.clone()

logs = pd.DataFrame()
logs[['loss_train', 'loss_val', 'ssim_train', 'ssim_val', 'mse_train', 'mse_val']] = None

model = UNet().to(device)
model.trainable_encoder(trainable=False)

loss_fn = nn.MSELoss()
optim = torch.optim.AdamW(model.parameters(), lr=lr / 25., weight_decay=0.02)
sched = torch.optim.lr_scheduler.OneCycleLR(optim, max_lr=lr, epochs=epochs, steps_per_
scaler = GradScaler()

```

```

In [ ]: best_ssim = -1e9
best_epoch = -1

print('training decoder only')

for epoch in tqdm(range(epochs)):

    model.train()

    if epoch == freeze_epochs:
        model.trainable_encoder(trainable=True)
        print('training encoder and decoder both')

    running_loss = 0.
    train_prog = tqdm(train_dl, total=len(train_dl))

    for img, mask in train_prog:

        with autocast():
            img, mask = img.to(device), mask.to(device)
            preds = model(img)

            loss = loss_fn(preds, mask)
            scaler.scale(loss).backward()
            scaler.unscale_(optim)
            nn.utils.clip_grad_norm_(model.parameters(), max_norm=2.0, norm_type=2)
            scaler.step(optim)
            scaler.update()
            sched.step()
            optim.zero_grad()

        running_loss += loss.item()

```

```

        train_prog.set_description(f' loss: {loss.item():.3f}')
        train_metrics(preds, mask)

    def img, mask, preds, loss

        m = train_metrics.compute()
        _ssim, _mse = m['StructuralSimilarityIndexMeasure'].cpu().item(), m['MeanSquaredError']
        logs.loc[epoch, ['loss_train', 'ssim_train', 'mse_train']] = (running_loss/len(train_dl), _ssim, _mse)
        train_metrics.reset()
        model.eval()

    with torch.no_grad():

        running_loss = 0.

        val_prog = tqdm(val_dl, total=len(val_dl))
        for img, mask in val_prog:

            with autocast():
                img, mask = img.to(device), mask.to(device)
                preds = model(img)
                loss = loss_fn(preds, mask)
                running_loss += loss.item()
            val_prog.set_description(f' loss: {loss.item():.3f}')

        val_metrics(preds, mask)

    def img, mask, preds, loss

        m = val_metrics.compute()
        _ssim, _mse = m['StructuralSimilarityIndexMeasure'].cpu().item(), m['MeanSquaredError']
        logs.loc[epoch, ['loss_val', 'ssim_val', 'mse_val']] = (running_loss/len(val_dl), _ssim, _mse)
        val_metrics.reset()

        if _ssim > best_ssim:
            best_ssim = _ssim
            best_epoch = epoch
            sd = model.state_dict()
            torch.save(sd, Path('./model-weights/nyu-v2-depth-resnext50_32x4d-unetplusplus+').with_suffix('.pt'))

    print(f"\n\n{logs.tail(1)}\n\n")

    with torch.no_grad():
        with autocast():
            img, mask = next(iter(test_dl))
            img, mask = img.to(device), mask.to(device)
            preds = model(img)
            plot_vals(
                img.cpu(),
                preds.cpu(),
                mask.cpu()
            )

    gc.collect()
    torch.cuda.empty_cache()

```

```

training decoder only
0%|          | 0/5 [00:00<?, ?it/s]
0%|          | 0/713 [00:00<?, ?it/s]
0%|          | 0/72 [00:00<?, ?it/s]

```

	loss_train	loss_val	ssim_train	ssim_val	mse_train	mse_val
0	0.03118	0.008758	0.578338	0.769319	0.031186	0.008767

image/target/prediction



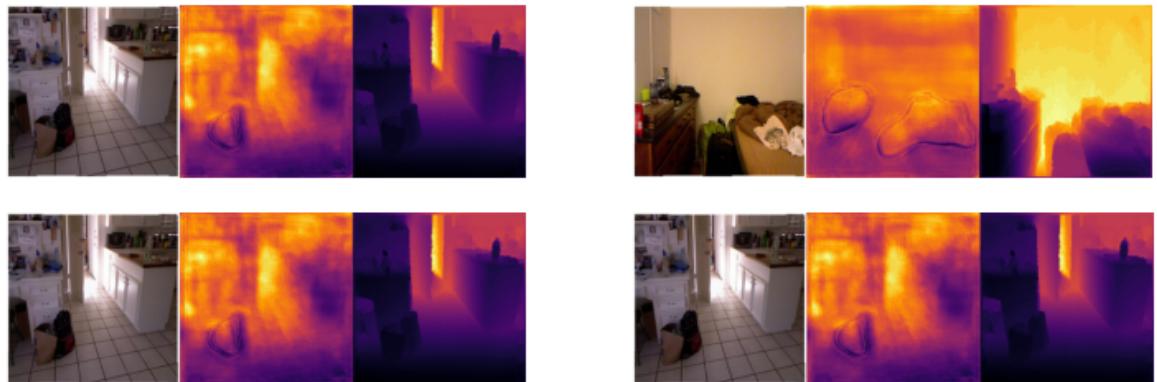
```

0% |           0/713 [00:00<?, ?it/s]
0% |           0/72 [00:00<?, ?it/s]

```

	loss_train	loss_val	ssim_train	ssim_val	mse_train	mse_val
1	0.010111	0.006337	0.853832	0.863473	0.010111	0.006352

image/target/prediction



training encoder and decoder both

```

0% |           0/713 [00:00<?, ?it/s]
0% |           0/72 [00:00<?, ?it/s]

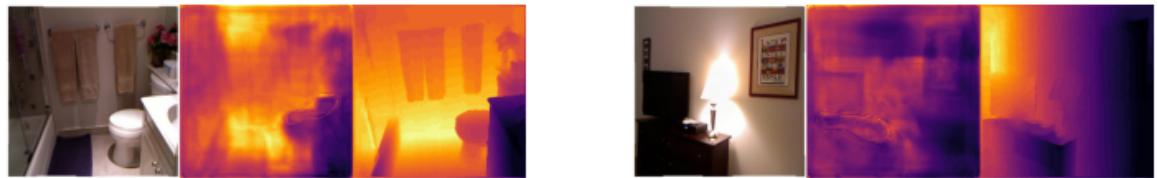
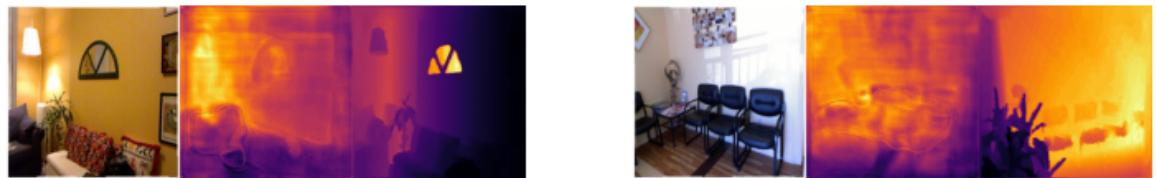
```

	loss_train	loss_val	ssim_train	ssim_val	mse_train	mse_val
2	0.010468	0.005191	0.875593	0.891569	0.010469	0.005197

**image/target/prediction**

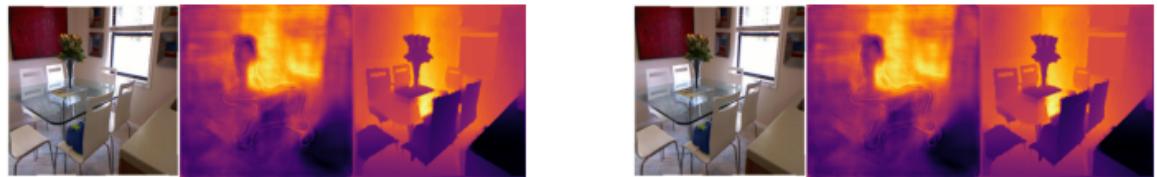
0% | 0/713 [00:00<?, ?it/s]  
0% | 0/72 [00:00<?, ?it/s]

loss_train	loss_val	ssim_train	ssim_val	mse_train	mse_val
3 0.006773	0.003099	0.898602	0.90524	0.006773	0.003106

**image/target/prediction**

0% | 0/713 [00:00<?, ?it/s]  
0% | 0/72 [00:00<?, ?it/s]

loss_train	loss_val	ssim_train	ssim_val	mse_train	mse_val
4 0.004926	0.002794	0.909537	0.908052	0.004927	0.002798

**image/target/prediction**

# Predictions

```
In [ ]: best_epoch
```

```
Out[ ]: 4
```

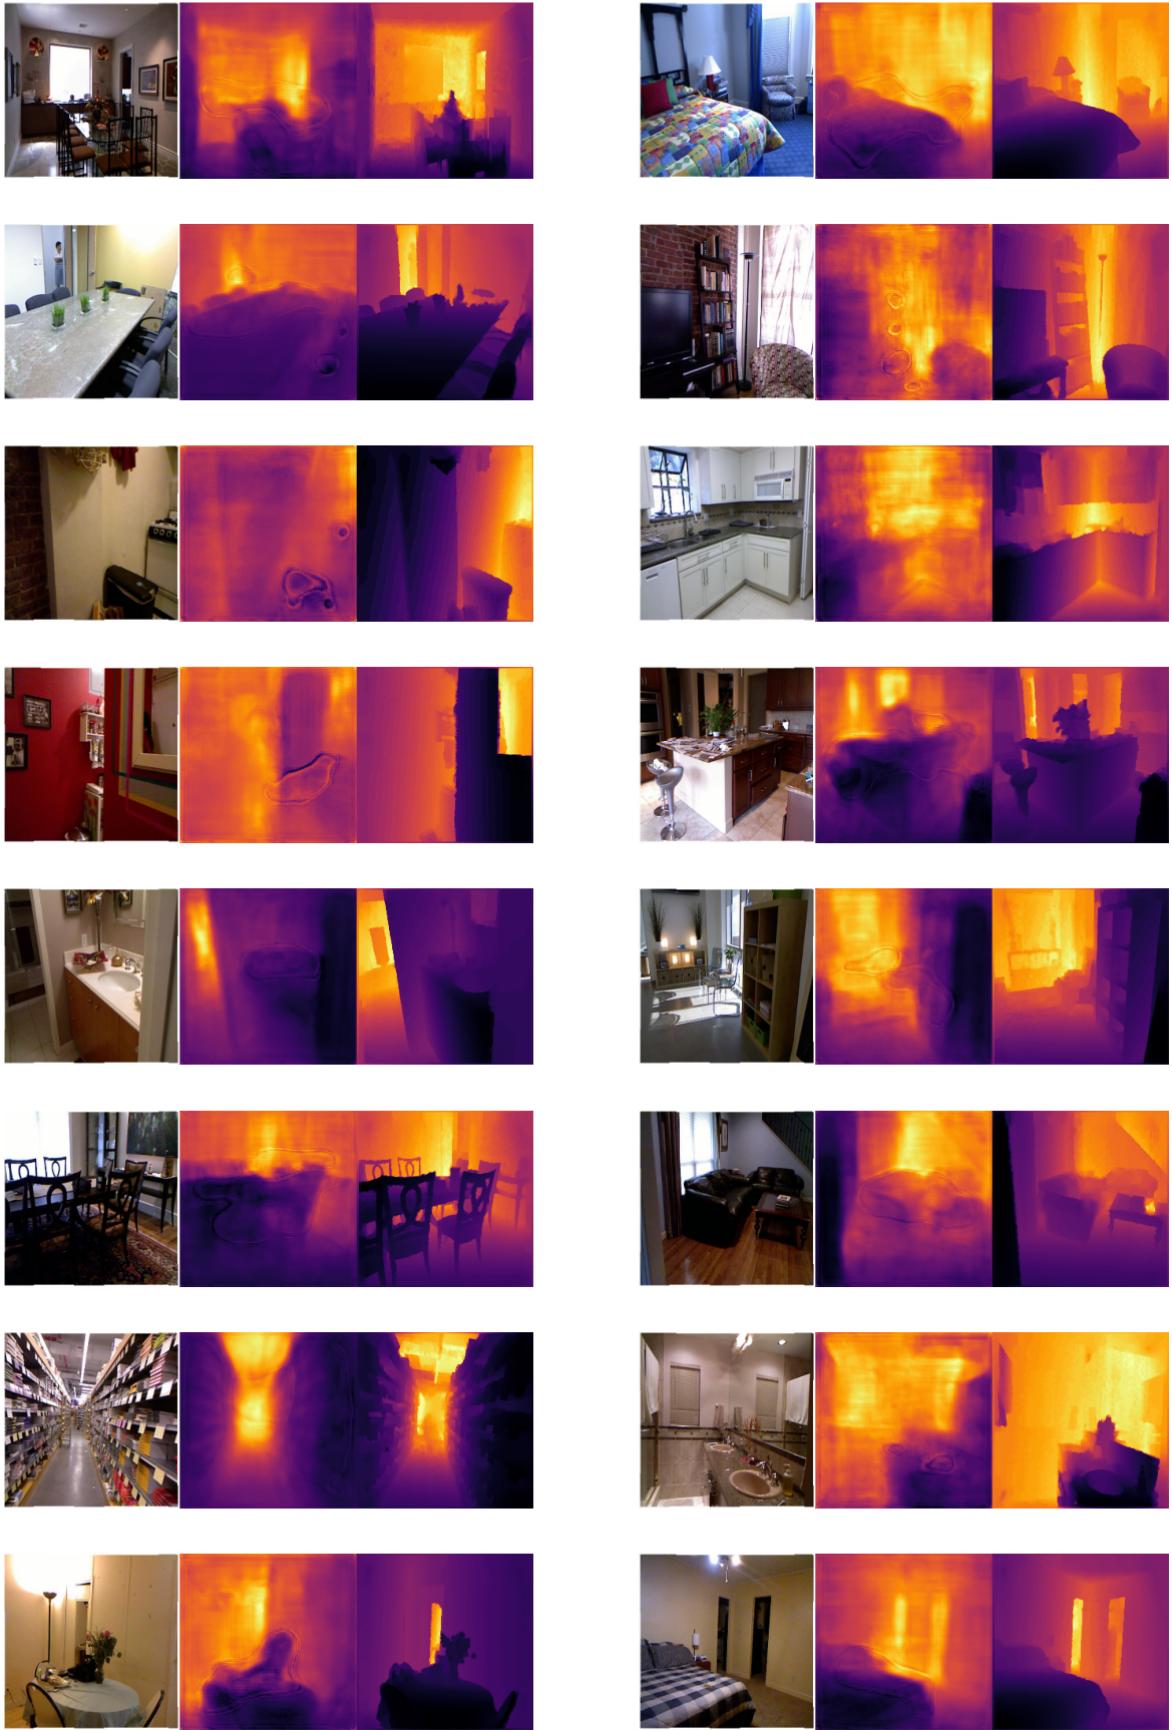
```
In [ ]: best_sd = torch.load(Path('./model-weights/nyu-v2-depth-resnext50_32x4d-unetplusplus'))  
model.load_state_dict(best_sd)
```

```
Out[ ]: <All keys matched successfully>
```

```
In [ ]: all_imgs, all_preds, all_targets = [], [], []  
with torch.no_grad():  
    with autocast():  
        for img, mask in tqdm(test_dl, total=len(test_dl)):  
            img, mask = img.to(device), mask.to(device)  
            preds = model(img)  
            all_imgs.append(img)  
            all_preds.append(preds)  
            all_targets.append(mask)  
  
test_metrics = metrics.clone()  
test_metrics(  
    torch.vstack(all_preds),  
    torch.vstack(all_targets)  
)  
m = test_metrics.compute()  
title = f"SSIM: {m['StructuralSimilarityIndexMeasure'].cpu().item():.3f} MSE: {m['Me  
plot_vals(  
    torch.vstack(all_imgs).cpu(),  
    torch.vstack(all_preds).cpu(),  
    torch.vstack(all_targets).cpu(),  
    n=16,  
    figsize=(10, 15),  
    title=title  
)
```

```
0% | 0/127 [00:00<?, ?it/s]
```

SSIM: 0.912 MSE: 0.003  
image/target/prediction

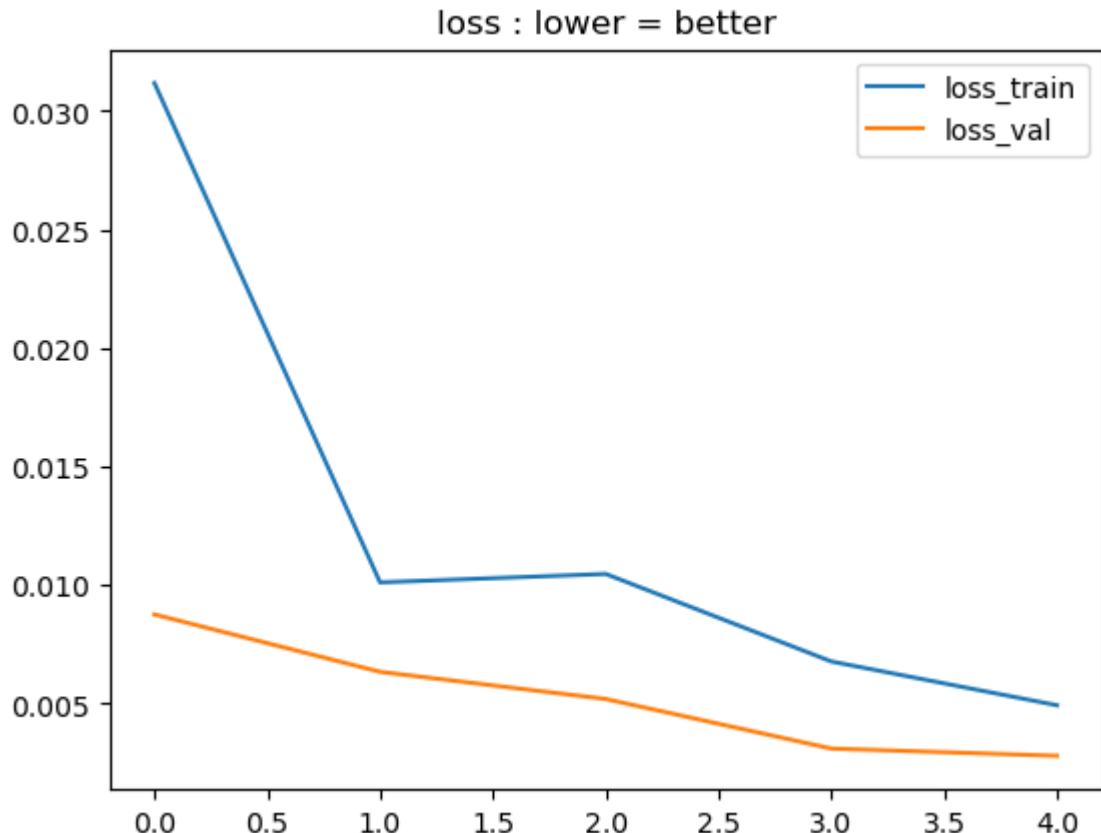


# Results

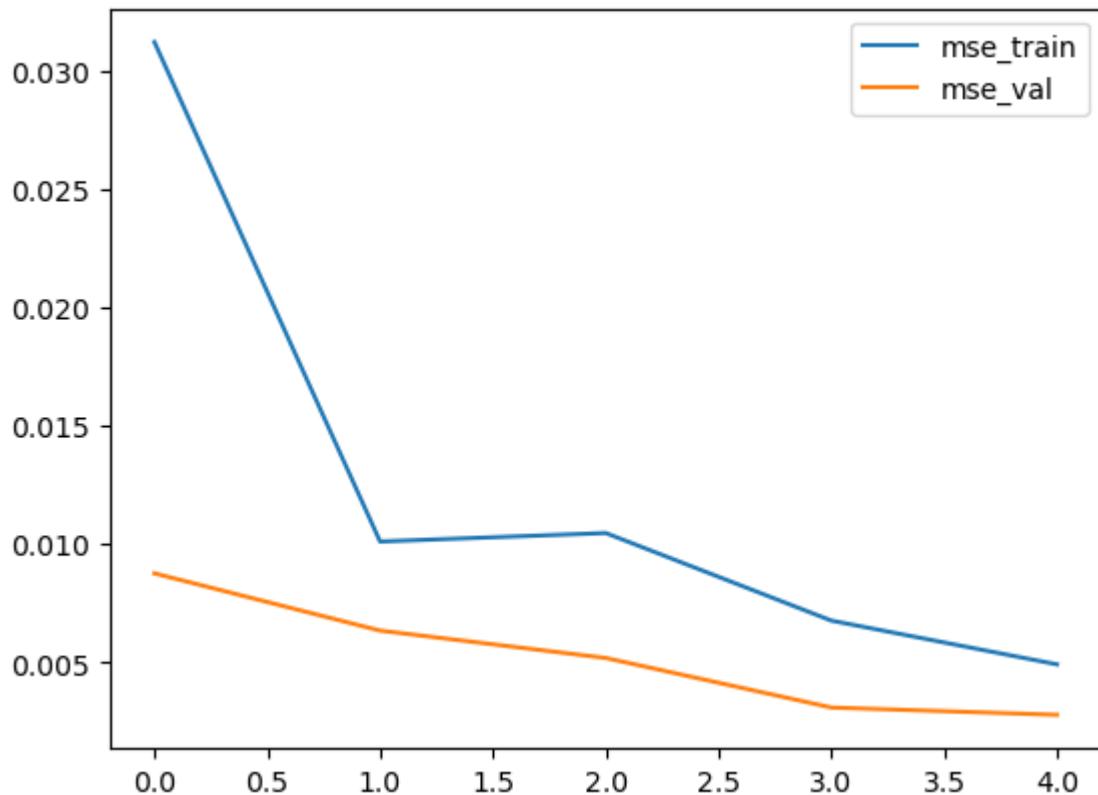
```
In [ ]: logs
```

```
Out[ ]:   loss_train  loss_val  ssim_train  ssim_val  mse_train  mse_val
0    0.03118  0.008758  0.578338  0.769319  0.031186  0.008767
1    0.010111  0.006337  0.853832  0.863473  0.010111  0.006352
2    0.010468  0.005191  0.875593  0.891569  0.010469  0.005197
3    0.006773  0.003099  0.898602  0.90524   0.006773  0.003106
4    0.004926  0.002794  0.909537  0.908052  0.004927  0.002798
```

```
In [ ]: logs['loss_train'].plot()
logs['loss_val'].plot()
plt.title('loss : lower = better')
plt.legend()
plt.show()
```



```
In [ ]: logs['mse_train'].plot()
logs['mse_val'].plot()
plt.title('mse : lower = better')
plt.legend()
plt.show()
```

**mse : lower = better**

```
In [ ]: logs['ssim_train'].plot()  
logs['ssim_val'].plot()  
plt.title('ssim : higher = better')  
plt.legend()  
plt.show()
```

**ssim : higher = better**